

Lecture 17

What we've done and what's to come

Announcements

- Final exam:
 - Wednesday 12/10, 8:30am-11:30am
 - Location: [see Ed post!]
- See website for more details, practice exams, etc!
 - “Official” practice exam posted by the end of this week
- See Ed for resources, study tips, etc!

More announcements

- Course feedback open!
 - A reminder should pop up on Canvas if you sign in!
 - Or fill it out directly here:
 - <http://course-evaluations.stanford.edu>.
 - Or go to Axess > My Academics > Course and Section Evaluations
- Your feedback is super-important!
- It will help us make the course better!

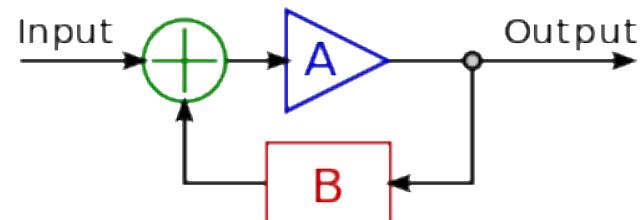


Figure 1: Feedback

More announcements

- This is the last class!
 - Time flies when you are having fun...
 - Or doing lots of algorithms problems... 😊
- Please keep in touch!
 - Stop by and say hi! (I live in CoDa W216 most of the time)

Help evaluate the Embedded Ethics Program!



- <https://tinyurl.com/EmbeddedEthicsFA25>
- 10-15 minute survey, taking it (or not) won't impact your grade in the class in any way, and teaching team won't know who participates.
- Option to provide your **Stanford email address** to receive a \$10 gift card, **up to the first 800 participants**. Compensation once per quarter (SUNet login required).
- Questions? Email embeddedethics@stanford.edu

Today

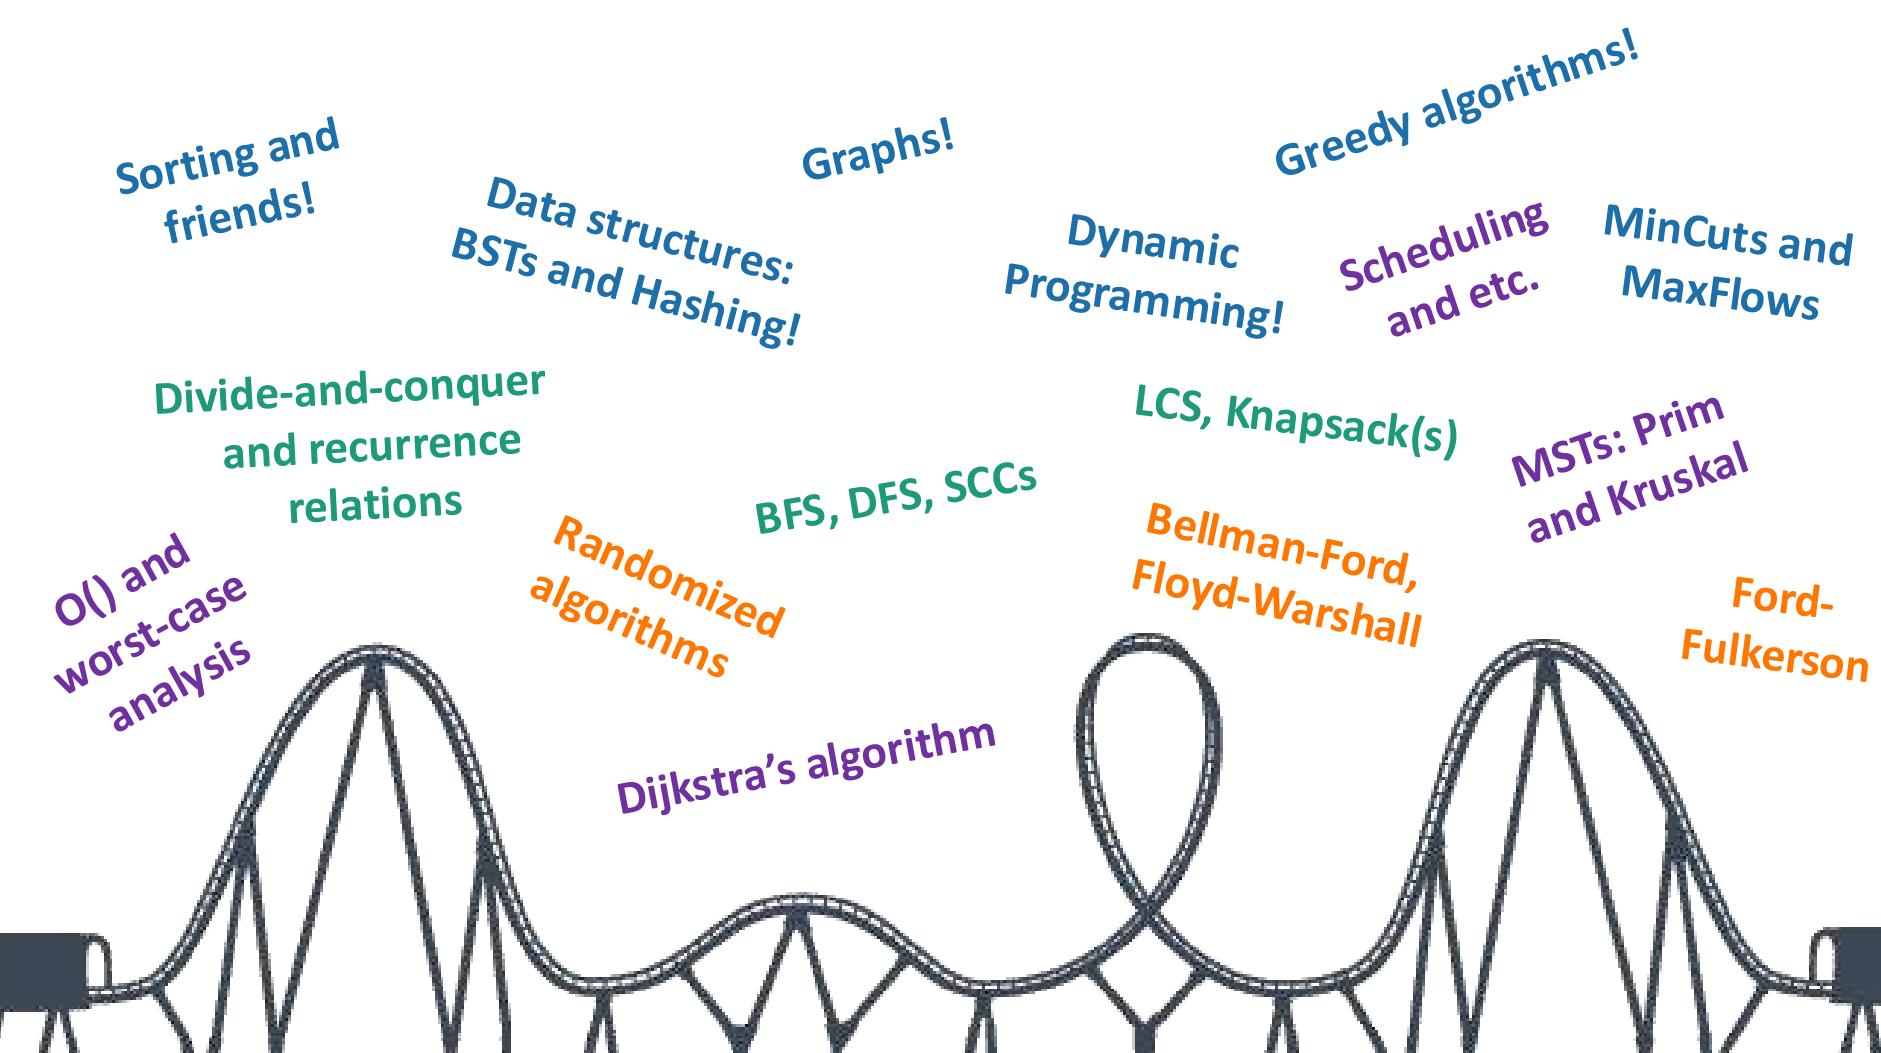
- What just happened?
 - A whirlwind tour of CS161



- What's next?
 - A few gems from future algorithms classes



It's been a fun ride...



What have we learned?

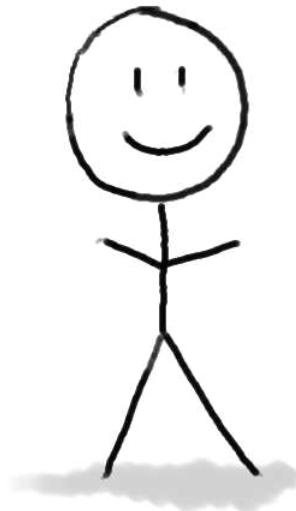
17 lectures in 13 slides.

General approach to algorithm design and analysis

Does it work?

Is it fast?

Can I do better?



Algorithm designer

To answer these questions we
need both **rigor** and **intuition**:



Plucky the
Pedantic Penguin

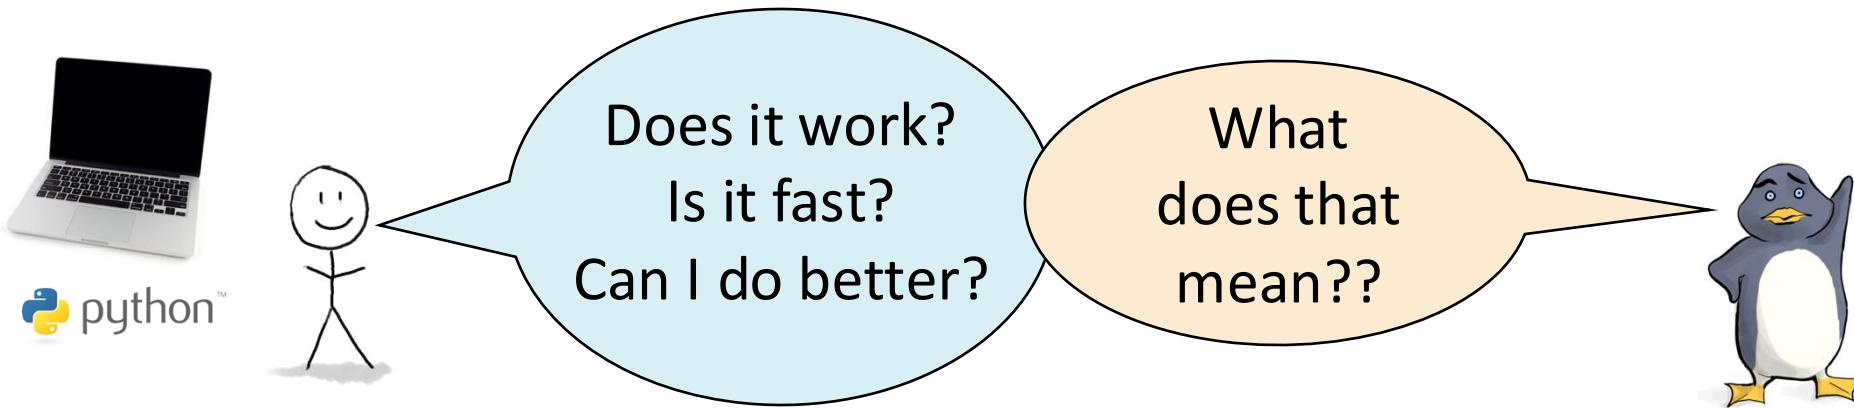
Detail-oriented
Precise
Rigorous



Lucky the
Lackadaisical Lemur

Big-picture
Intuitive
Hand-wavey

We needed more details

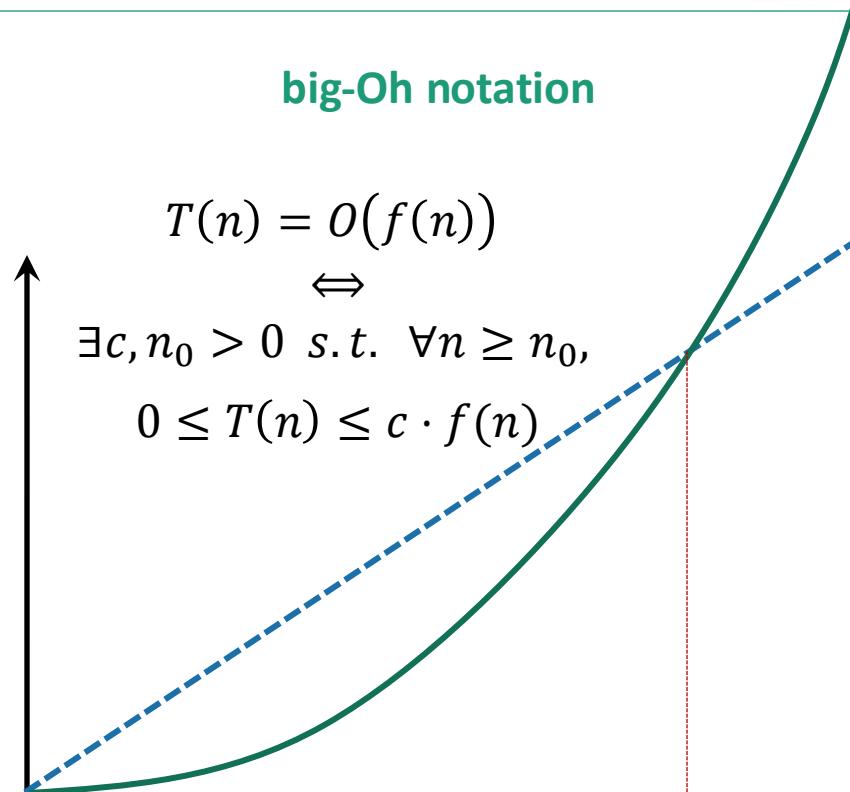


Worst-case analysis



big-Oh notation

$$T(n) = O(f(n)) \iff \exists c, n_0 > 0 \text{ s.t. } \forall n \geq n_0, 0 \leq T(n) \leq c \cdot f(n)$$



Algorithm design paradigm: divide and conquer

- Like MergeSort!
- Or Karatsuba's algorithm!
- Or SELECT!
- How do we analyze these?

By careful
analysis!

Tree method
Substitution Method

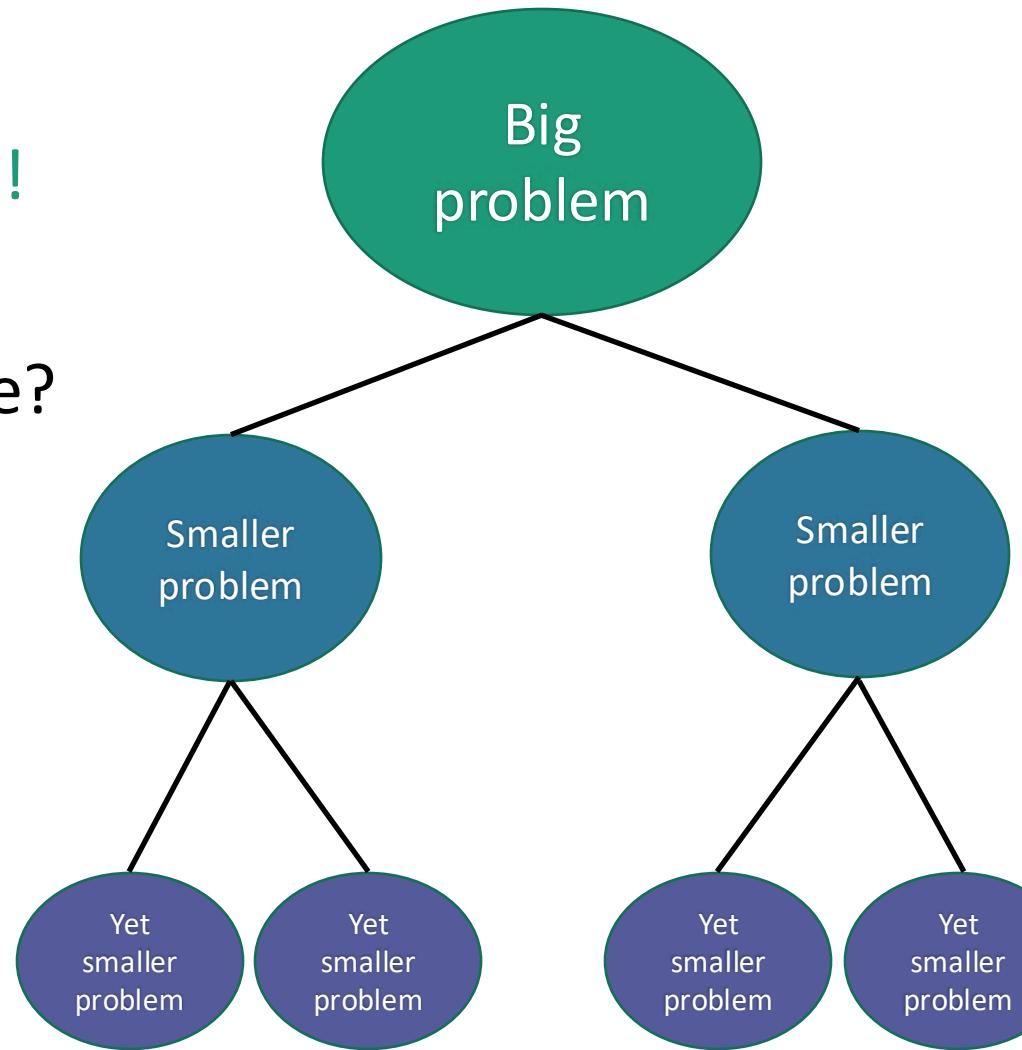


Plucky the
Pedantic Penguin

Useful shortcut, the
master method is.



Jedi master Yoda



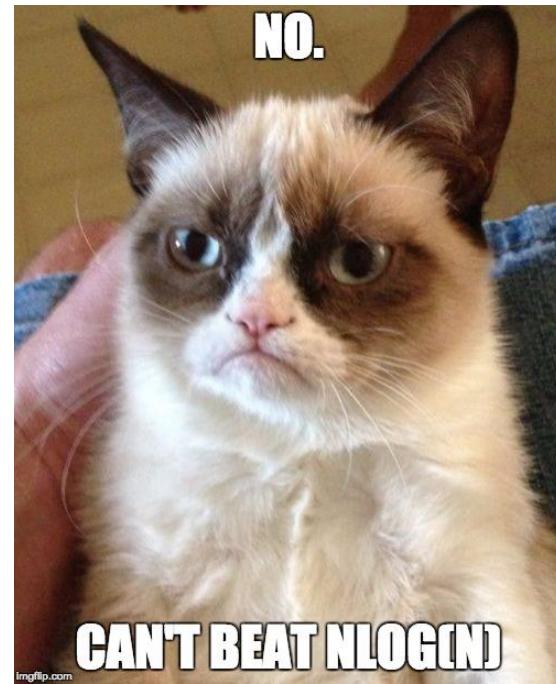
While we're on the topic of sorting Why not use randomness?

- We analyzed **QuickSort!**
- Still worst-case input, but we use randomness after the input is chosen.
- Always correct, usually fast.
 - This is a Las Vegas algorithm



All this sorting is making me wonder... Can we do better?

- Depends on who you ask:



- **RadixSort** takes time $O(n)$ if the objects are, for example, small integers!
- Can't do better in a **comparison-based** model.



beyond sorted arrays/linked lists: Binary Search Trees!

- Useful data structure!
- Especially the self-balancing ones!

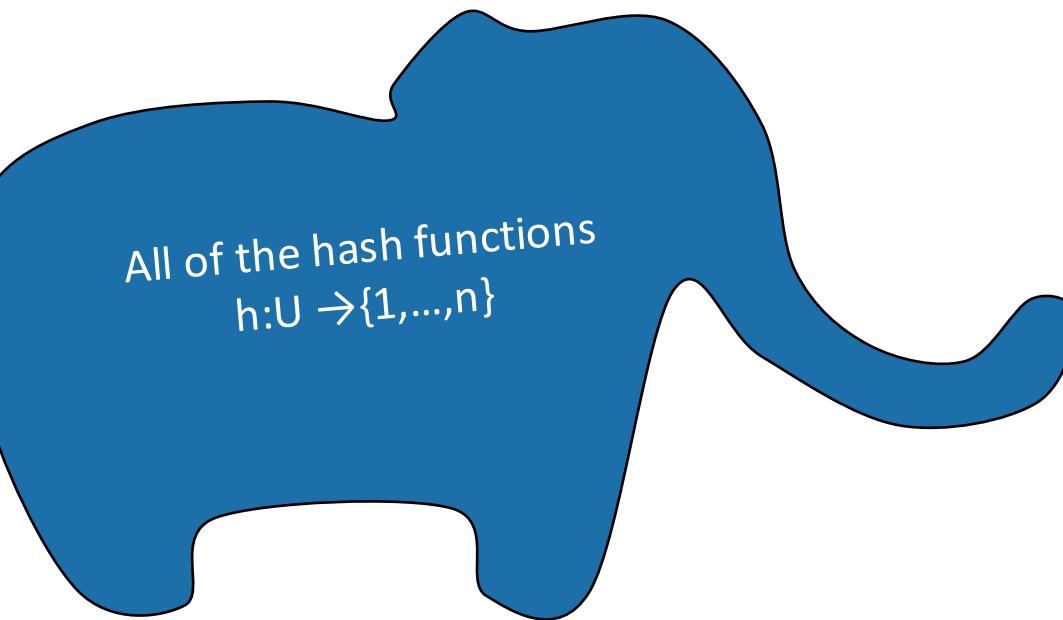
Red-Black tree!

Maintain balance by stipulating that
black nodes are balanced, and
that there aren't too many **red
nodes**.

It's just good sense!



Another way to store things Hash tables!



Choose h randomly from a universal hash family.



It's better if the hash family is small!
Then it takes less space to store h .

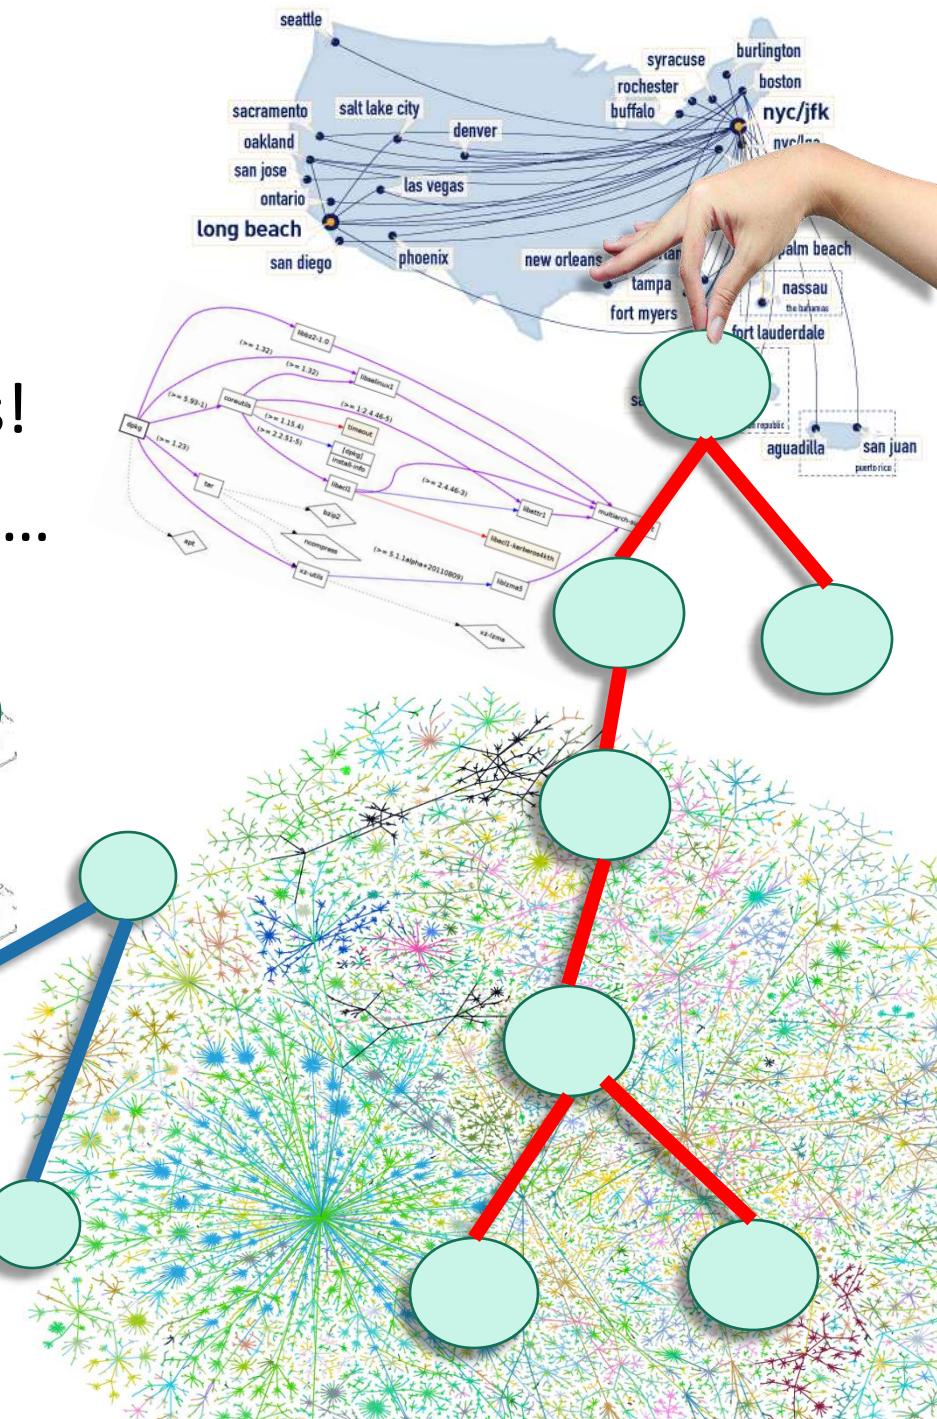
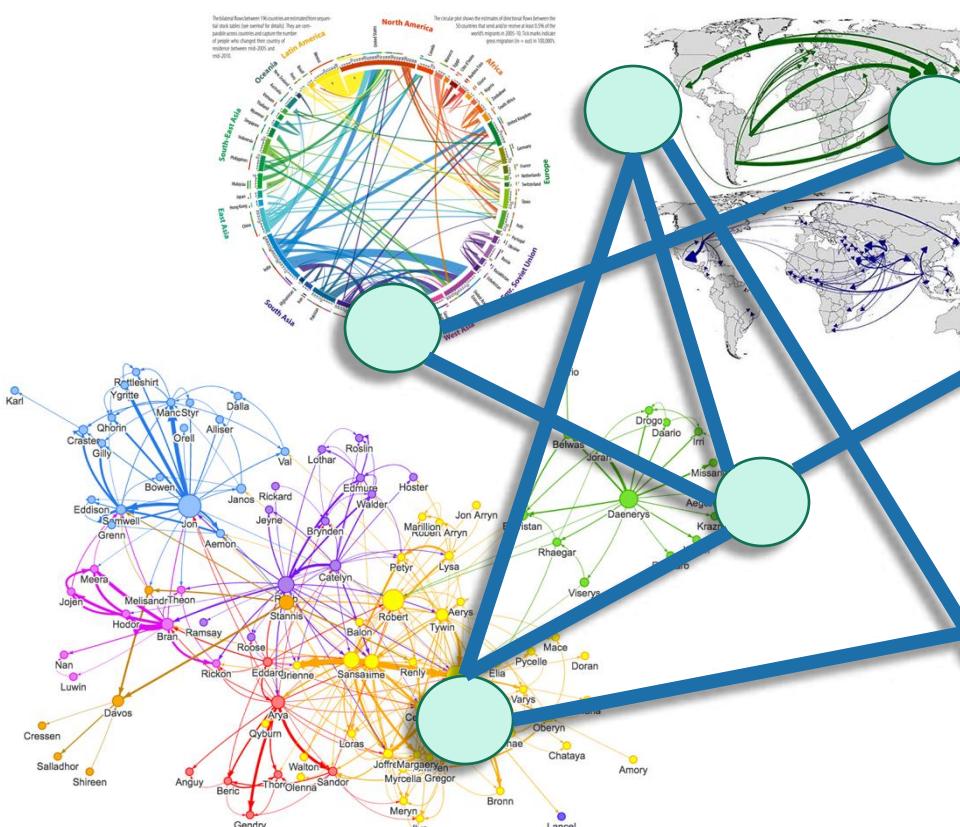


hash function h



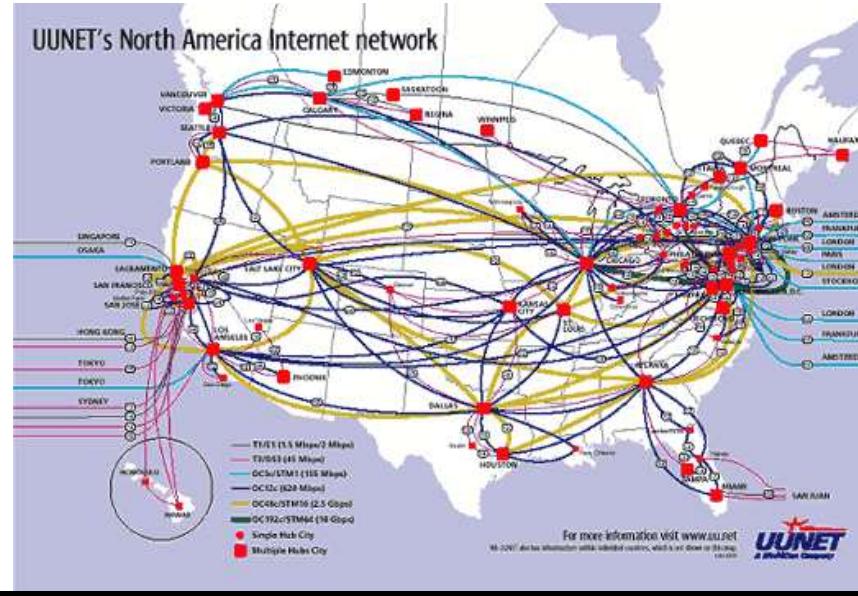
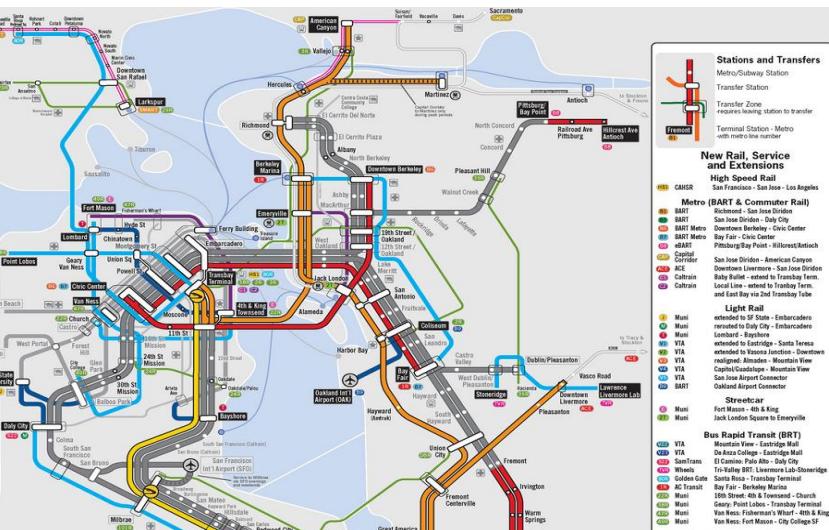
OMG GRAPHS

- BFS, DFS, and applications!
 - SCCs, Topological sorting, ...



A fundamental graph problem: shortest paths

- Eg, transit planning, packet routing, ...
- Dijkstra!
- Bellman-Ford!
- Floyd-Warshall!



```
[DN0a22a0e3:~ mary$ traceroute -a www.ethz.ch
traceroute to www.ethz.ch (129.132.19.216), 64 hops max, 52 byte packets
1 [AS0] 10.34.160.2 (10.34.160.2) 38.168 ms 31.272 ms 28.841 ms
2 [AS0] cwa-vrtr.sunet (10.21.196.28) 33.769 ms 28.245 ms 24.373 ms
3 [AS32] 171.66.2.229 (171.66.2.229) 24.468 ms 20.115 ms 23.223 ms
4 [AS32] hpr-svl-rtr-vlan8.sunet (171.64.255.235) 24.644 ms 24.962 ms 17.453 ms
5 [AS2152] hpr-svl-hpr2--stan-ge.cenic.net (137.164.27.161) 22.129 ms 4.902 ms 3.642 ms
6 [AS2152] hpr-lax-hpr3--svl-hpr3-100ge.cenic.net (137.164.25.73) 12.125 ms 43.361 ms 32.3 ms
7 [AS2152] hpr-i2-lax-hpr2-r-e.cenic.net (137.164.26.201) 40.174 ms 38.399 ms 34.499 ms
8 [AS0] et-4-0-0.4079.sdn-sw.lasv.net.internet2.edu (162.252.70.28) 46.573 ms 23.926 ms
9 [AS0] et-5-1-0.4079.rtsw.salt.net.internet2.edu (162.252.70.31) 30.424 ms 25.770 ms 23.1 ms
10 [AS0] et-4-0-0.4079.sdn-sw.denv.net.internet2.edu (162.252.70.8) 47.454 ms 57.273 ms 73.1 ms
11 [AS0] et-4-1-0.4079.rtsw.kans.net.internet2.edu (162.252.70.11) 70.825 ms 67.809 ms 62.1 ms
12 [AS0] et-4-1-0.4070.rtsw.chic.net.internet2.edu (198.71.47.206) 77.937 ms 57.421 ms 63.6 ms
13 [AS0] et-0-1-0.4079.sdn-sw.ashb.net.internet2.edu (162.252.70.60) 77.682 ms 71.993 ms 73.1 ms
14 [AS0] et-4-1-0.4079.rtsw.wash.net.internet2.edu (162.252.70.65) 71.565 ms 74.988 ms 71.0 ms
15 [AS21320] internet2-gw.mx1.lon.uk.geant.net (62.40.124.44) 154.926 ms 145.606 ms 145.872 ms
16 [AS21320] ae0.mx1.lon.uk.geant.net (62.40.98.79) 146.565 ms 146.604 ms 146.801 ms
17 [AS21320] ae0.mx1.par.fr.geant.net (62.40.98.77) 153.289 ms 184.995 ms 152.682 ms
18 [AS21320] ae2.mx1.gen.ch.geant.net (62.40.98.153) 160.283 ms 160.104 ms 164.147 ms
19 [AS21320] swicel1-100ge-0-3-0-1.switch.ch (62.40.124.22) 162.068 ms 160.595 ms 163.095 ms
20 [AS559] swizh1-100ge-0-1-0-1.switch.ch (130.59.36.94) 165.824 ms 164.216 ms 163.983 ms
21 [AS559] swiez3-100ge-0-1-0-4.switch.ch (130.59.38.109) 164.269 ms 164.370 ms 163.929 ms
22 [AS559] rou-gw-lee-tengig-to-switch.ethz.ch (192.33.92.1) 164.082 ms 170.645 ms 165.372 ms
23 [AS559] rou-fw-rz-rz-gw.ethz.ch (192.33.92.169) 164.773 ms 165.193 ms 172.158 ms
```

Bellman-Ford and Floyd-Warshall
were examples of...

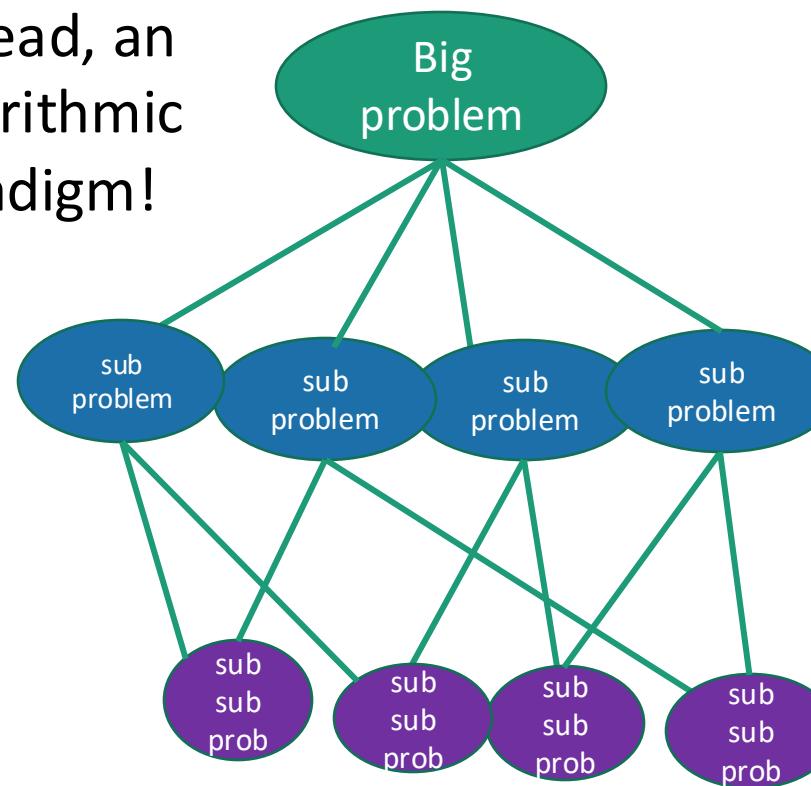
Dynamic Programming!

- Not programming in an action movie.



- Step 1:** Identify optimal substructure.
- Step 2:** Find a recursive formulation for the value of the optimal solution.
- Steps 3-5:** Use dynamic programming: fill in a table to find the answer!

Instead, an
algorithmic
paradigm!



We saw many other examples, including Longest Common Subsequence and Knapsack Problems.

Sometimes we can take even better advantage of optimal substructure...with

Greedy algorithms

- Make a series of choices, and commit!



- Intuitively we want to show that our greedy choices never rule out success.
- Rigorously, we usually analyzed these by induction.
- Examples!

- Activity Selection
- Job Scheduling
- Huffman Coding
- Minimum Spanning Trees

Prim's algorithm:
greedily grow a tree

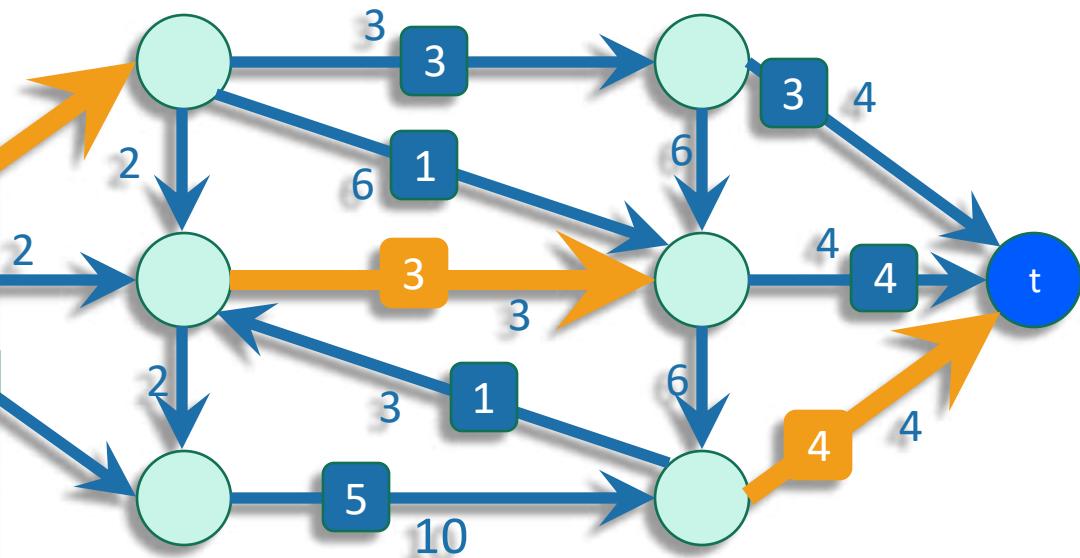
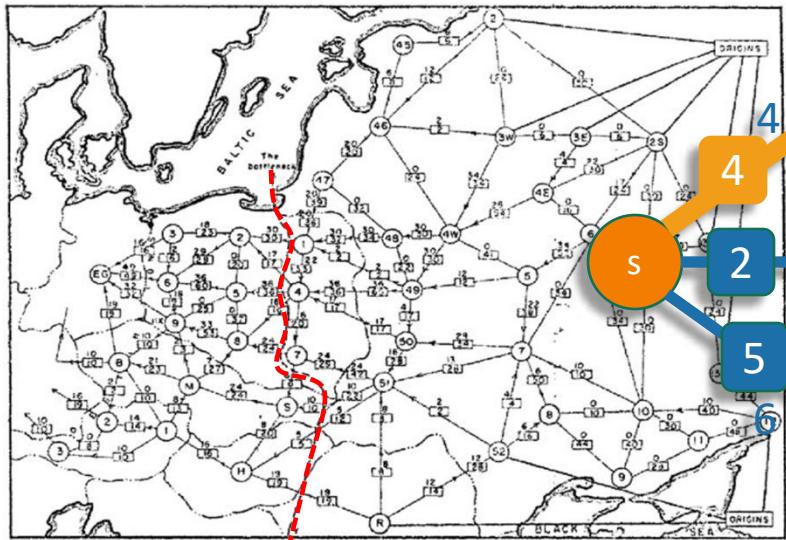


Kruskal's algorithm:
greedily grow a forest



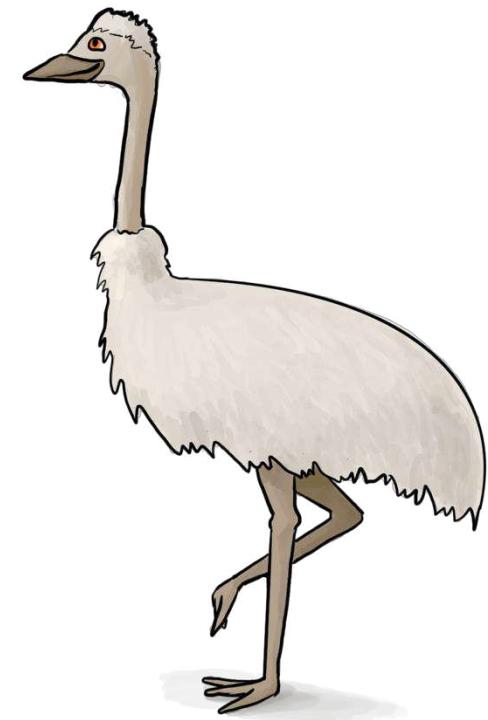
Cuts and flows

- minimum s-t cut is the same as maximum s-t flow!
- Ford-Fulkerson can find them!
 - Increase flow along augmenting paths!
- Applications:
 - routing supplies or interrupting supply chains
 - Giving ice cream to Stanford students



Along the way

- Embedded EthiCS!
 - Idealization, Abstraction, Measurement
 - Wicked problems and benign problems
- If/when we are developing algorithms in real life, we now have some things to watch out for and think about!



Emery the Ethical Emu

And now we're here



What have we learned?

- A few algorithm design paradigms:
 - Divide and conquer, Dynamic Programming, Greedy
- A few analysis tools:
 - Worst-case analysis, asymptotic analysis, recurrence relations, probability tricks, proofs by induction
- A few common objects:
 - Graphs, arrays, trees, hash functions
- A LOT of examples!



What have we learned? We've filled out a toolbox

- Tons of examples give us intuition about what algorithmic techniques might work when.
- The technical skills make sure our intuition works out.



But there's lots more out there



- What's next???

Want more classes?

findSomeTheoryCourses():

- go to theory.stanford.edu
- Click on “People”
- Look at what we’re teaching!

- CS154 – Introduction to Complexity
- CS166 – Data Structures
- CS168 – The Modern Algorithmic Toolbox
- MS&E 212 – Combinatorial Optimization
- CS250 – Error Correcting Codes
- CS254 – Computational Complexity
- CS255 – Introduction to Cryptography
- CS261 – Optimization and Algorithmic Paradigms
- CS264 – Beyond Worst-Case Analysis
- CS265 – Randomized Algorithms
- CS269Q – Quantum Computing
- CS269O – Introduction to Optimization Theory
- EE364A/B – Convex Optimization I and II

*...and many many more
upper-level topics courses!*

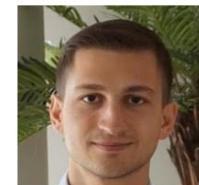
Faculty



Nima Anari



John C. Mitchell



Vasilis Syrgkanis



Dan Boneh



Omer Reingold



Li-Yang Tan



Adam Bouland



Aviad Rubinstein



Gregory Valiant



Moses Charikar



Amin Saberi



Ellen Vitercik



Ashish Goel



Tselil Schramm



Jan Vondrak



Tengyu Ma



Aaron Sidford



Mary Wootters

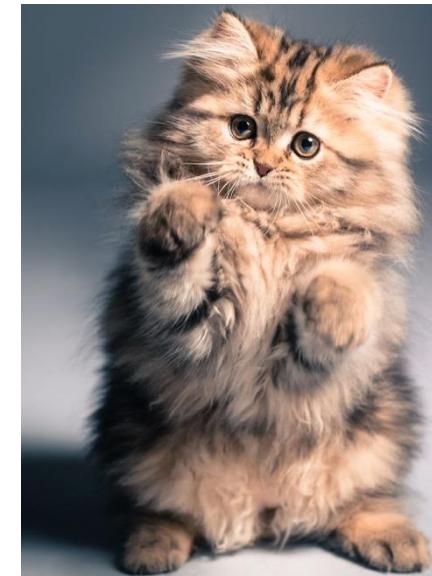
Today

A few gems

- Linear programming
- Random projections
- Low-degree polynomials



This will be pretty fluffy,
without much detail –
take more CS theory
classes for more detail!



**NOTHING AFTER THIS POINT
WILL BE ON THE FINAL EXAM**

Linear Programming

- This is a fancy name for optimizing a linear function subject to linear constraints.
- For example:

Maximize

$$x + y$$

subject to

$$x \geq 0$$

$$y \geq 0$$

$$4x + y \leq 2$$

$$x + 2y \leq 1$$

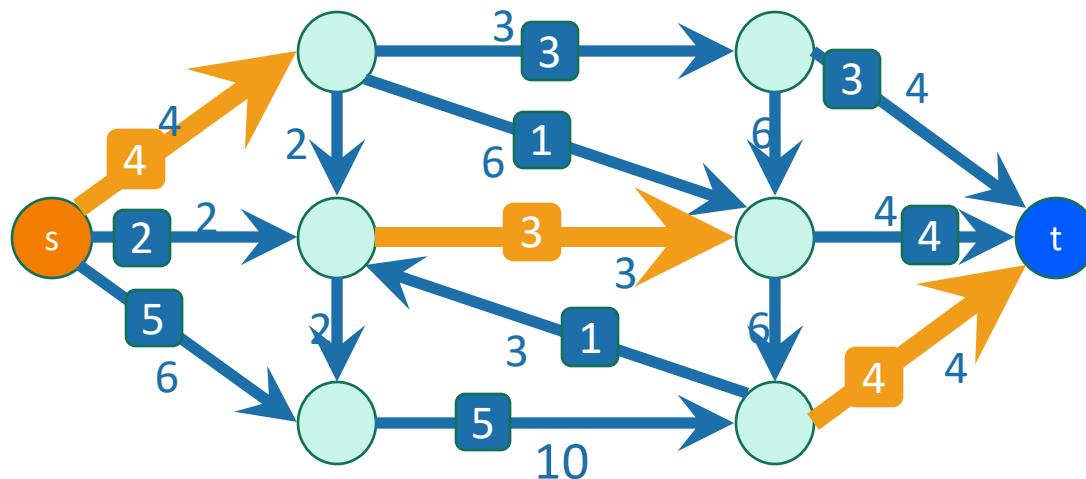
- It turns out the be an extremely general problem.

Actually we just saw it!

Maximize
the sum of the
flows leaving s

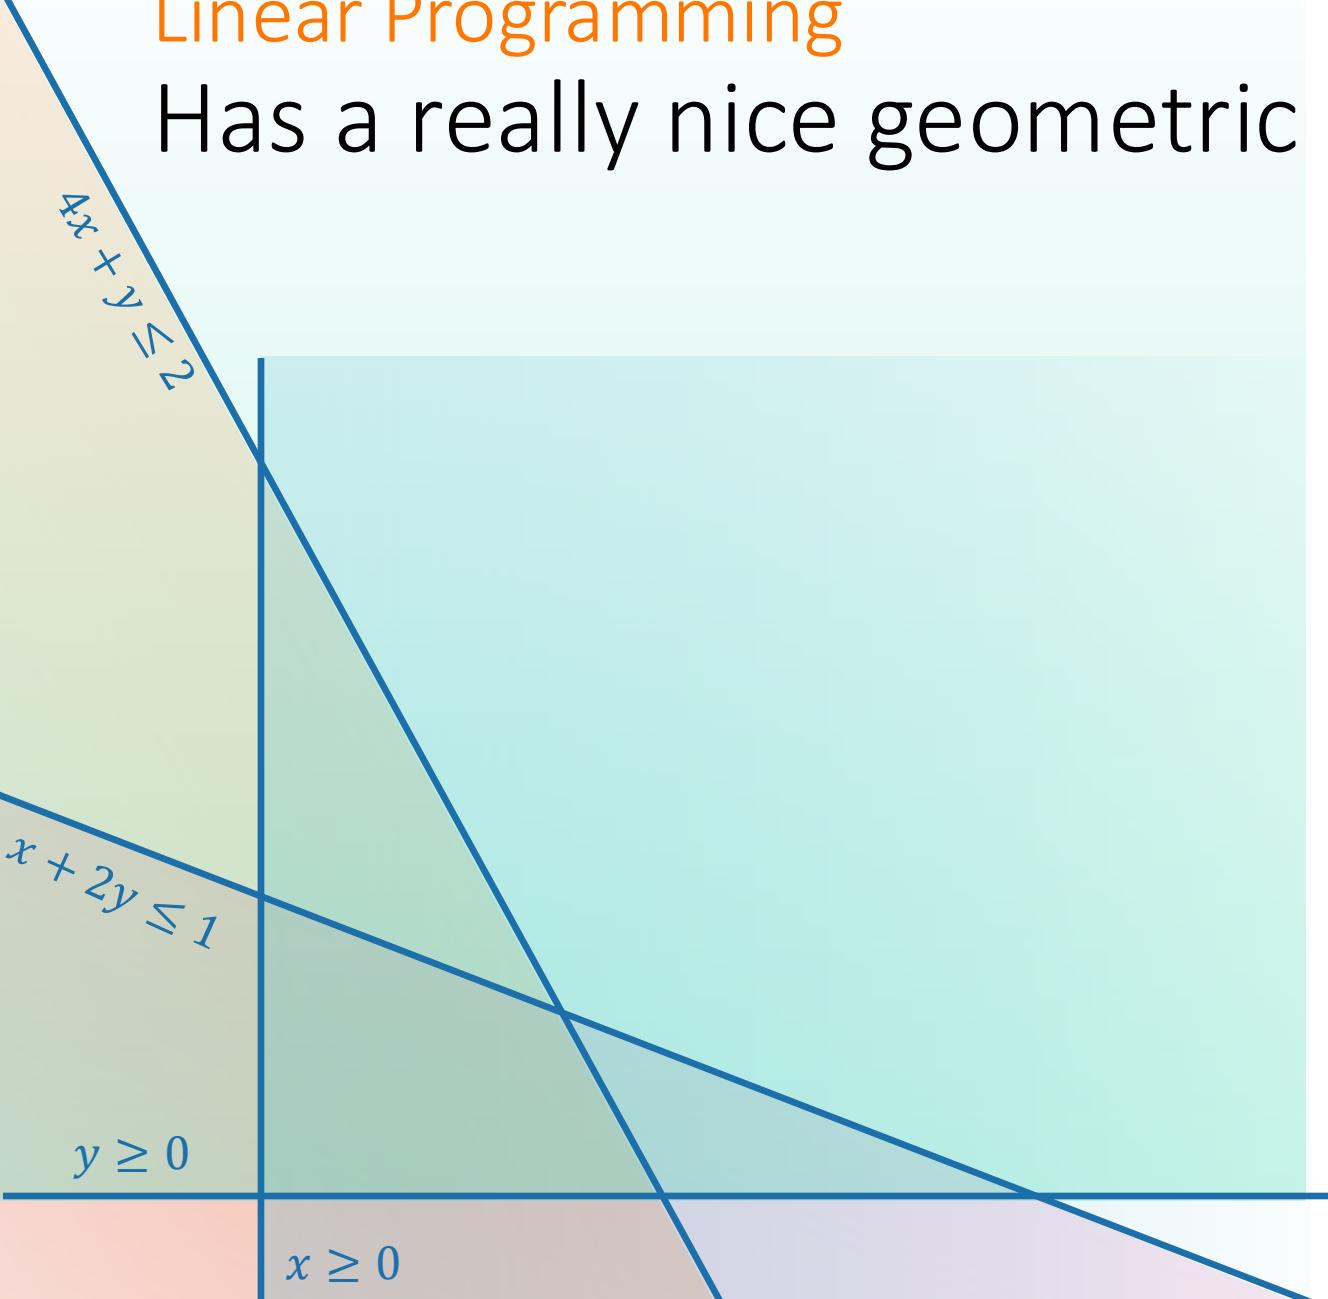
subject to

- None of the flows are bigger than the edge capacities
- At every vertex, stuff going in = stuff going out.



Linear Programming

Has a really nice geometric intuition



Maximize

$$x + y$$

subject to

$$x \geq 0$$

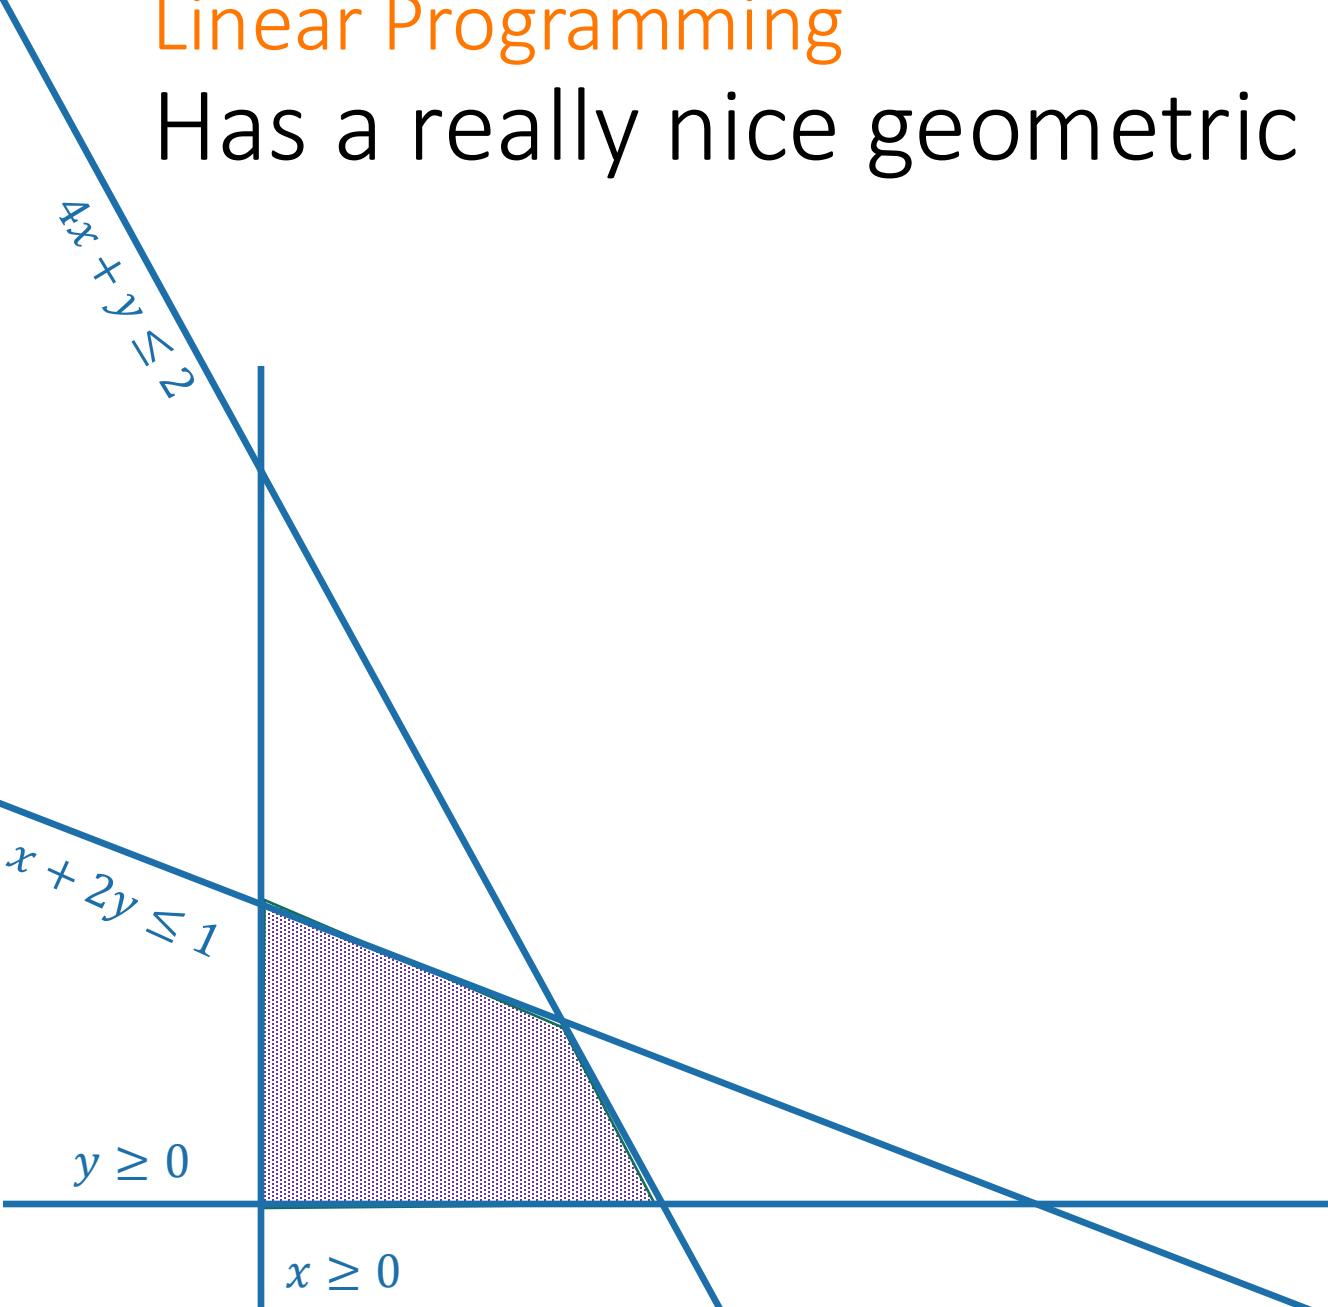
$$y \geq 0$$

$$4x + y \leq 2$$

$$x + 2y \leq 1$$

Linear Programming

Has a really nice geometric intuition



Maximize

$$x + y$$

subject to

$$x \geq 0$$

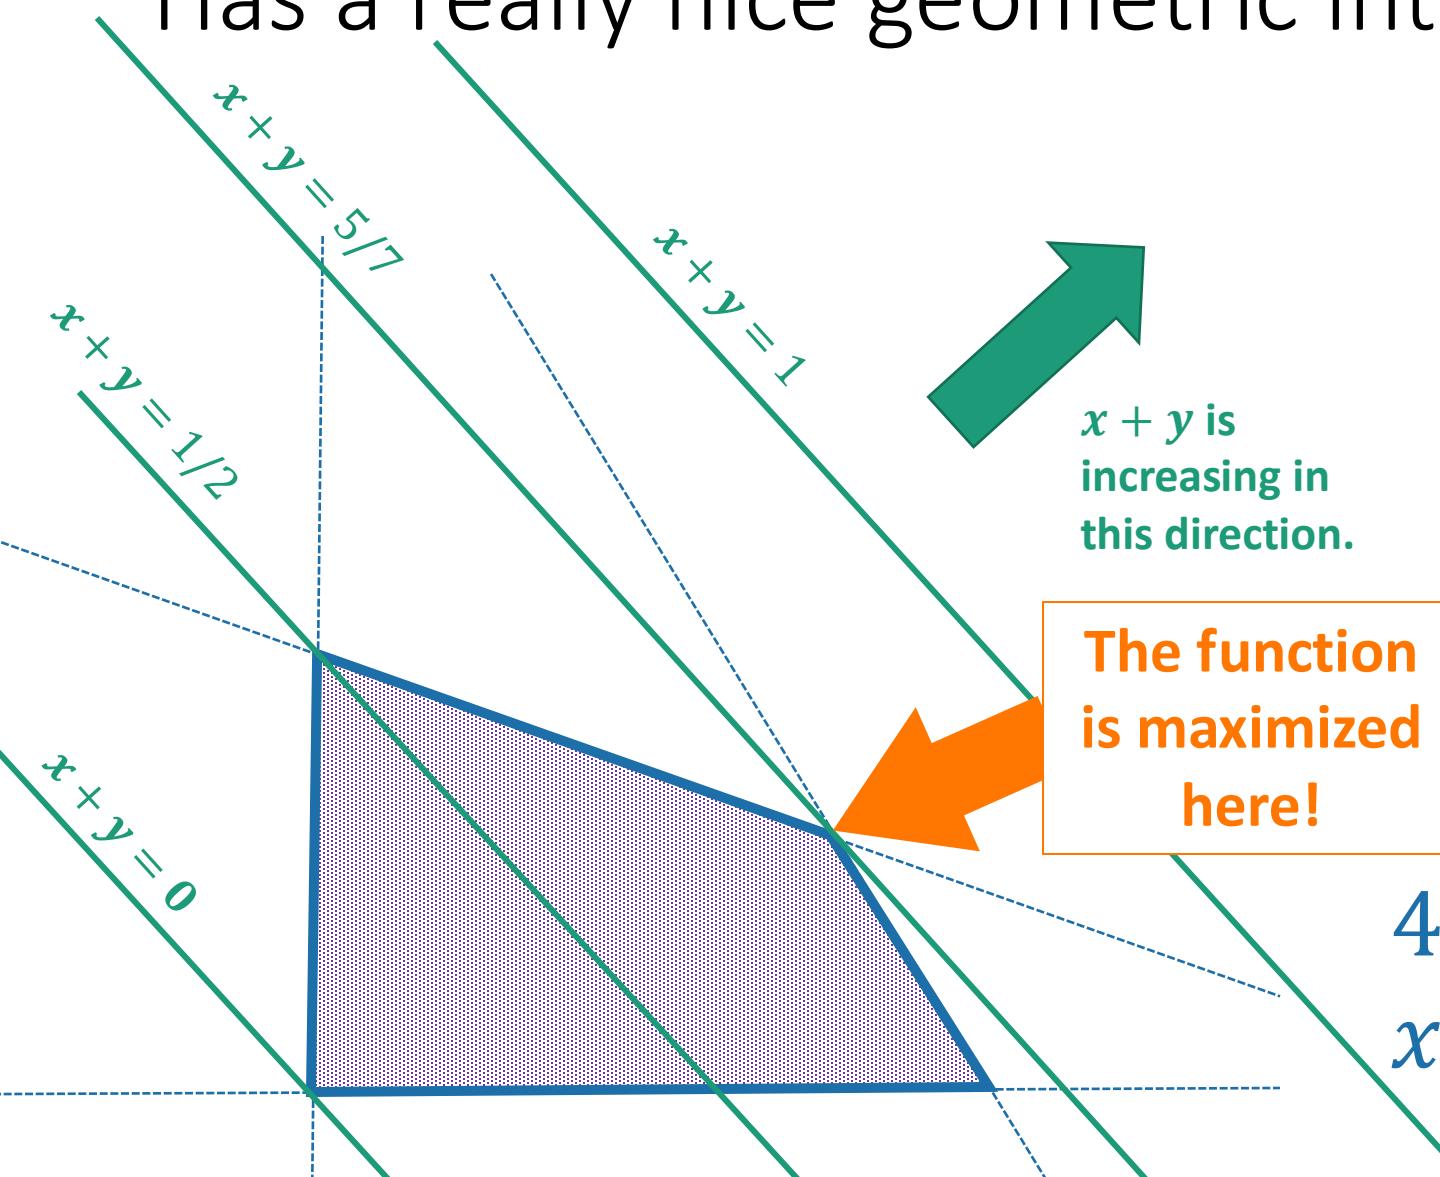
$$y \geq 0$$

$$4x + y \leq 2$$

$$x + 2y \leq 1$$

Linear Programming

Has a really nice geometric intuition



Maximize

$$x + y$$

subject to

$$x \geq 0$$

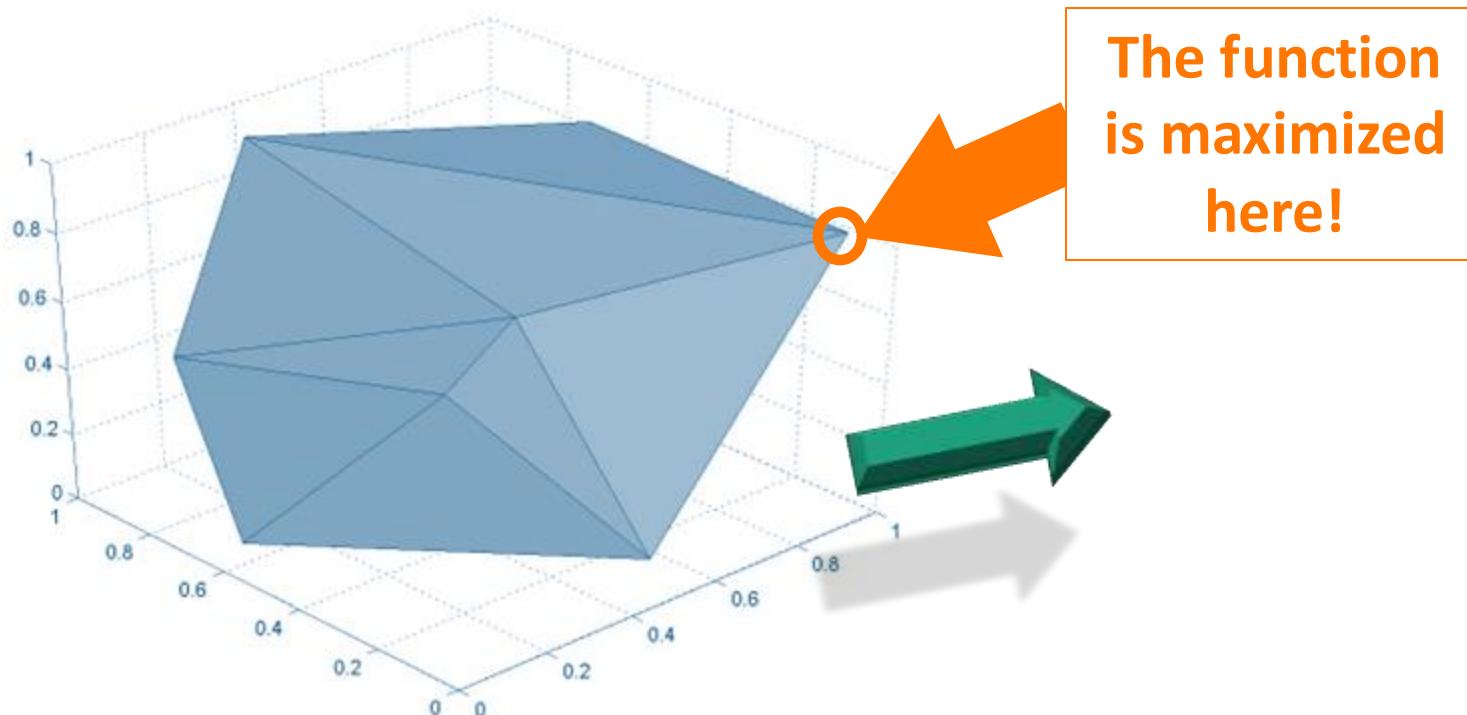
$$y \geq 0$$

$$4x + y \leq 2$$

$$x + 2y \leq 1$$

In general

- The constraints define a **polytope**
- The function defines a **direction**
- We just want to find the vertex that is **furthest** in that direction.



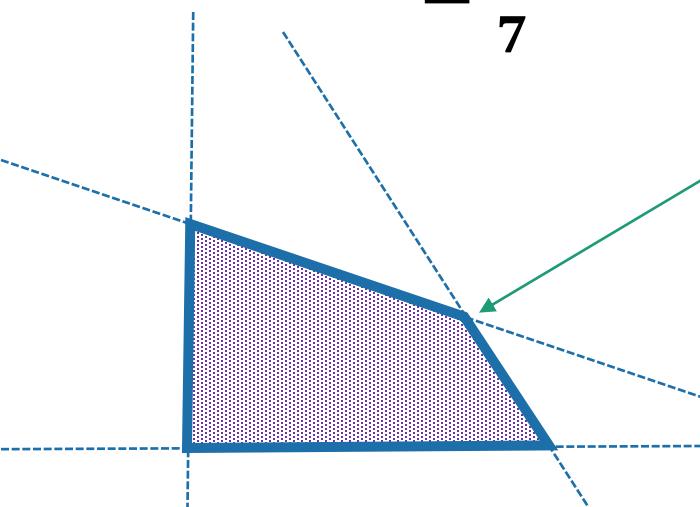
Duality

How do we know we have an optimal solution?

I claim that the optimum is $5/7$.

Proof: say x and y satisfy the constraints.

- $x + y = \frac{1}{7}(4x + y) + \frac{3}{7}(x + 2y)$
- $\leq \frac{1}{7} \cdot 2 + \frac{3}{7} \cdot 1$
- $= \frac{5}{7}$



You can check this point has value $5/7$...but how would we prove it's optimal other than by eyeballing it?

Maximize
 $x + y$

subject to

$$x \geq 0$$

$$y \geq 0$$

$$4x + y \leq 2$$

$$x + 2y \leq 1$$

cute, but

How did you come up with $1/7, 3/7$?

I claim that the optimum is $5/7$.

Proof: say x and y satisfy the constraints.

- $x + y \leq (4x + y) + (x + 2y)$
- $\leq \cdot 2 + 1$
- $=$
- I want to choose things to put here
- So that I minimize this
- Subject to these things

Maximize

$$x + y$$

subject to

$$x \geq 0$$

$$y \geq 0$$

$$4x + y \leq 2$$

$$x + 2y \leq 1$$

Note: it's not immediately obvious how to turn that into a linear program, this is just meant to convince you that it's plausible.

In this case the dual is:
 $\min 2w + z$ s.t. $w, z \geq 0$,
 $4w + z \geq 1$ and $w + 2z \geq 1$

That's a linear program!

- How did I find those special values $1/7, 3/7$?
- I solved some linear program. Minimize the upper bound you get, subject to the proof working.
- It's called the **dual program**.

Maximize stuff
subject to stuff

Primal

The optimal values are
the same!

Minimize other stuff
subject to other stuff

Dual

We've actually already seen this too

The Min-Cut Max-Flow Theorem!

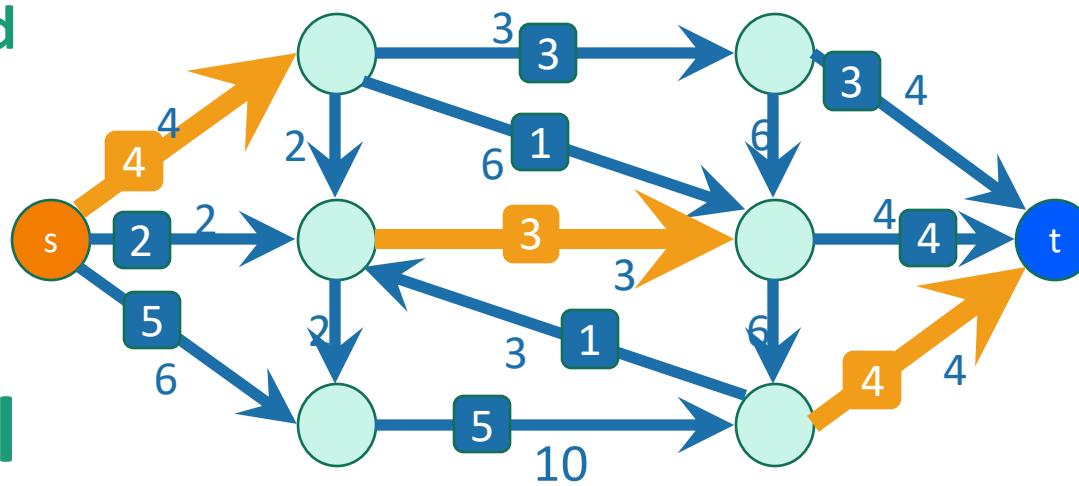
Maximize the sum of the flows leaving s
s.t.
All the flow constraints are satisfied

The optimal values are the same!

Minimize the sum of the capacities on a cut
s.t.
it's a legit cut

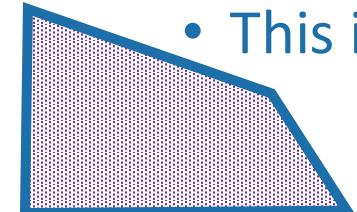
Primal

Dual



LPs and Duality are really powerful

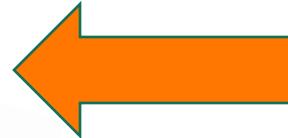
- This **general phenomenon** shows up all over the place
 - Min-Cut Max-Flow is a special case.
- Duality helps us reason about an optimization problem
 - The dual provides a **certificate** that we've solved the primal.
 - eg, if you have a cut and a flow with the same value, you must have found a max flow and a min cut.
- We can solve LPs quickly!
 - For example, by intelligently bouncing around the vertices of the feasible region.
 - This is an **extremely powerful algorithmic primitive**.



Today

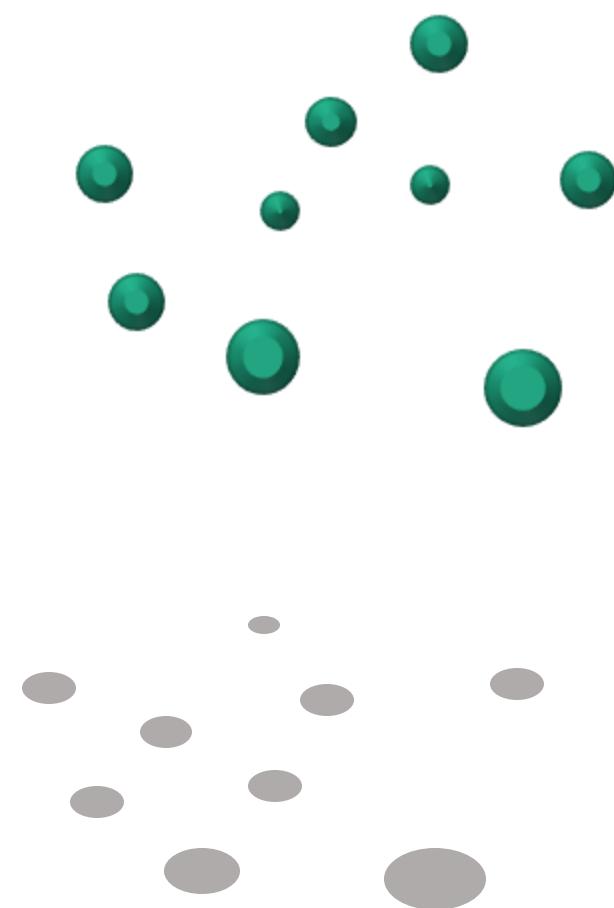
A few gems

- Linear programming
- Random projections
- Low-degree polynomials



One of my favorite tricks

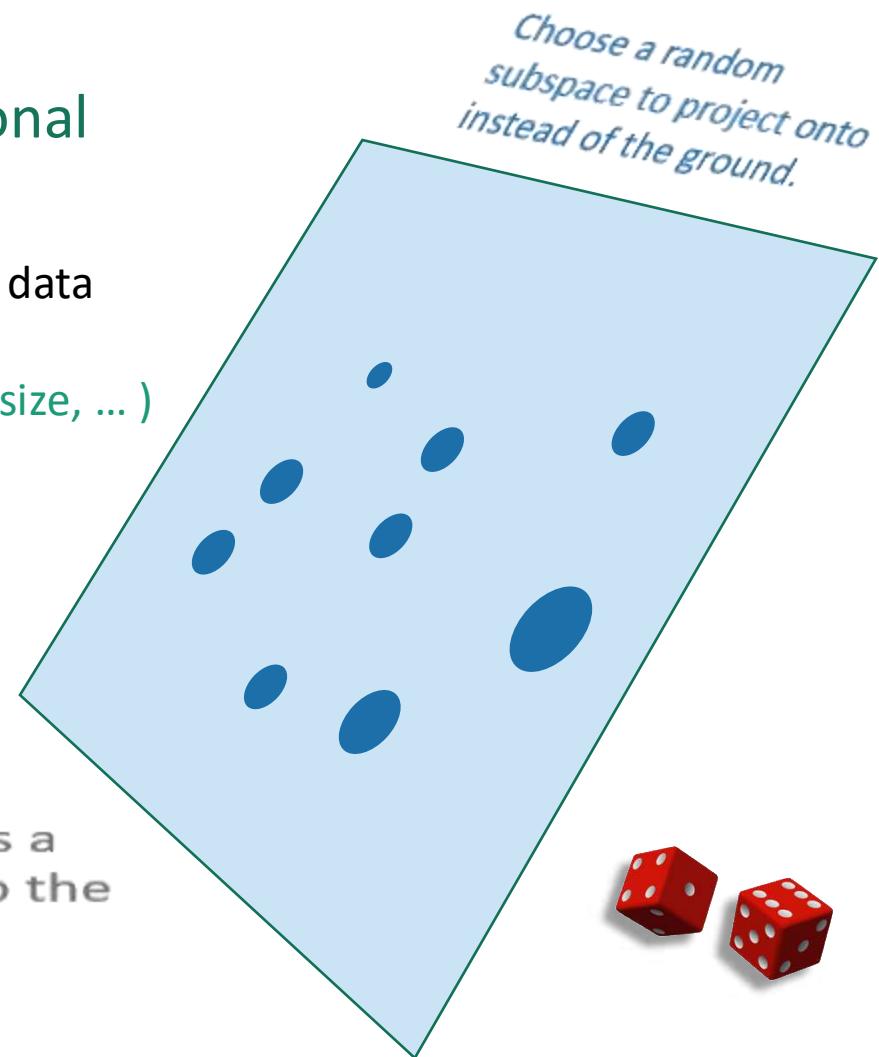
Take a random projection and hope for the best.



High-dimensional
set of points

For example, each data
point is a vector
(age, height, shoe size, ...)

Their shadow is a
projection onto the
ground.

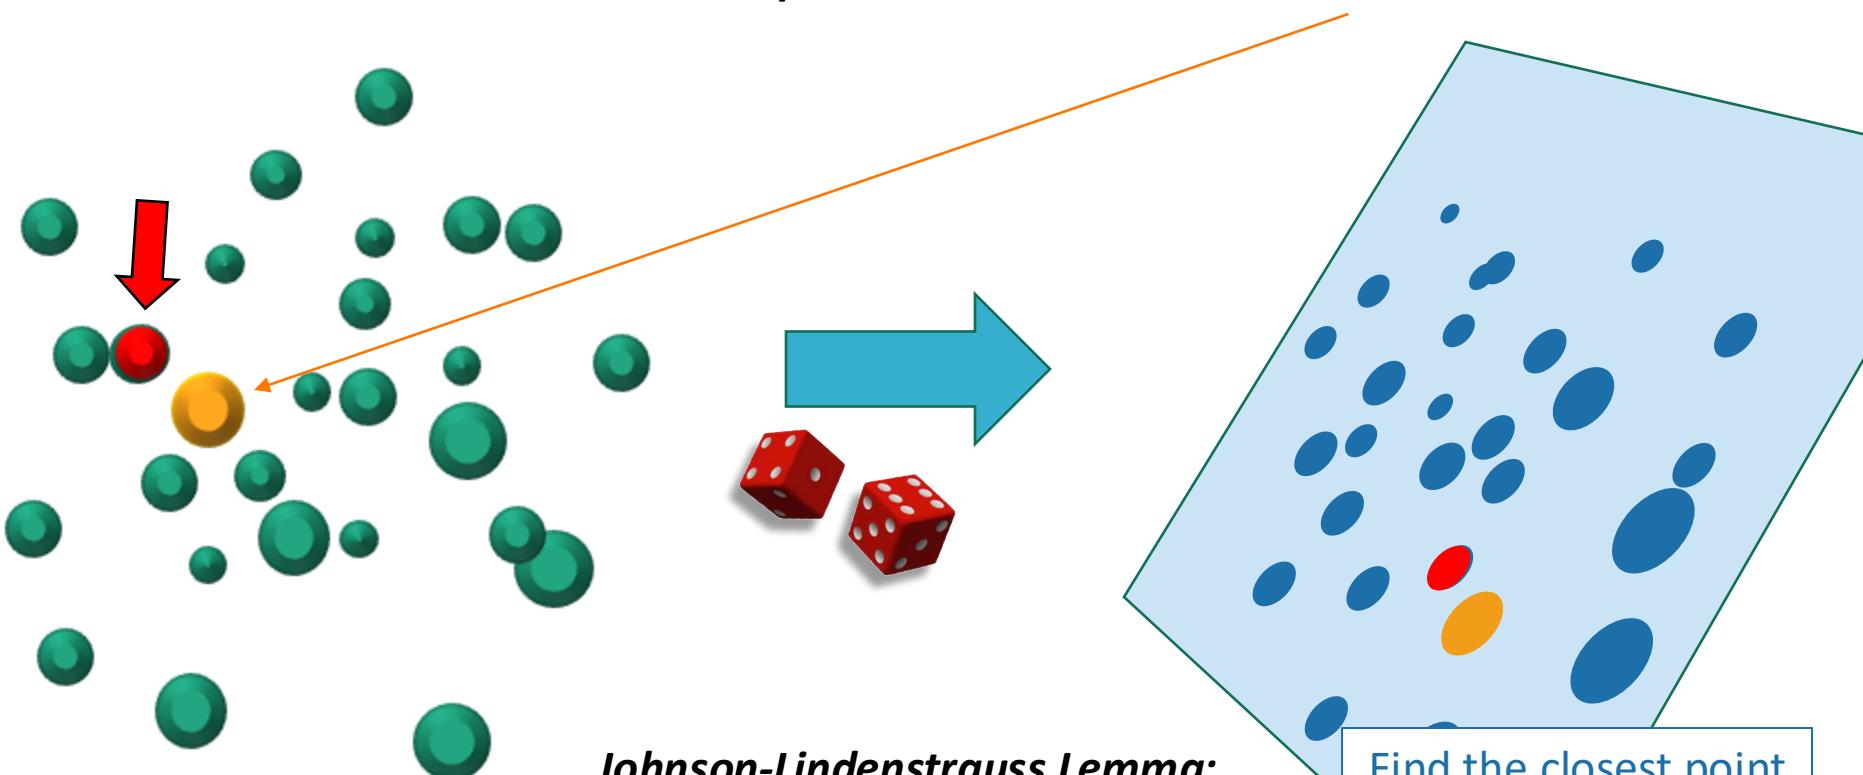


Why would we do this?

- High dimensional data takes a long time to process.
- Low dimensional data can be processed quickly.
- “**THEOREM**”: Random projections approximately preserve properties of data that you care about.

Example: nearest neighbors

- I want to find which point is closest to **this one**.



That takes a really long time in high dimensions.

Johnson-Lindenstrauss Lemma:
Euclidean distance is approximately preserved by random projections.

Find the closest point down here, you're probably pretty correct.

Another example: Compressed Sensing

- Start with a sparse vector
 - Mostly zero or close to zero

(5, 0, 0, 0, 0, 0.01, 0.01, 5.8, 32, 14, 0, 0, 0, 12, 0, 0, 5, 0, .03)

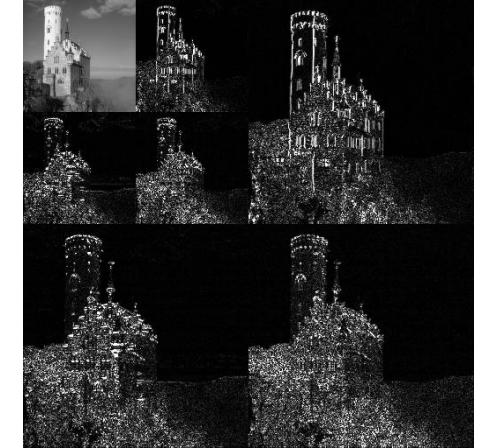
- For example:



This image is sparse

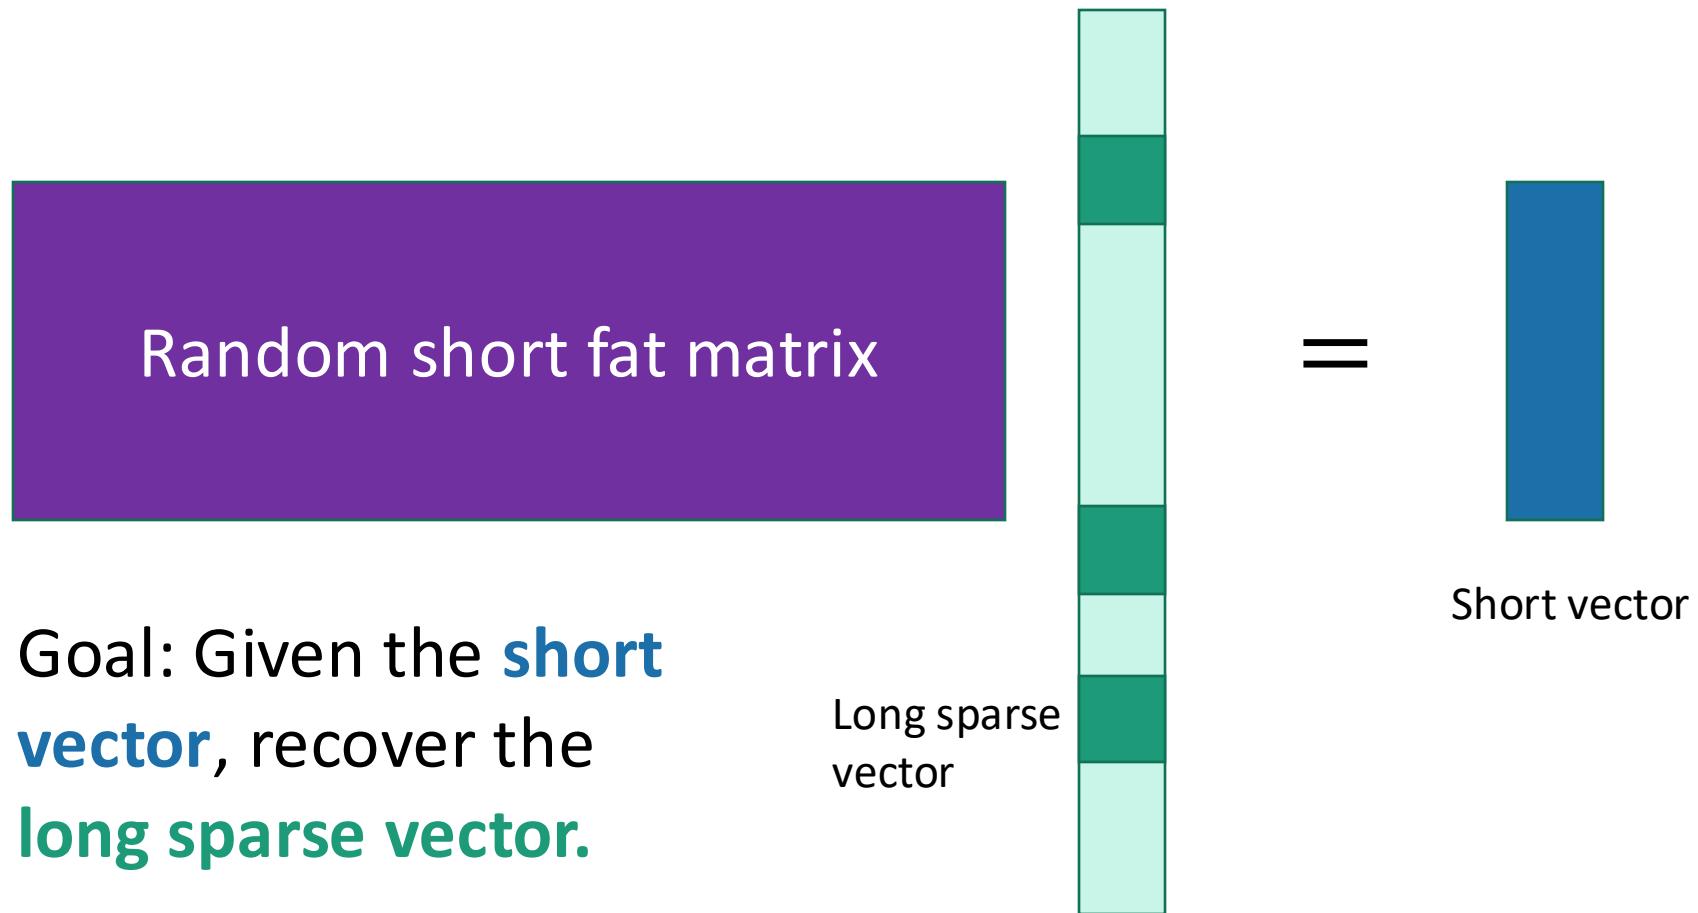


This image is sparse after I
take a wavelet transform.



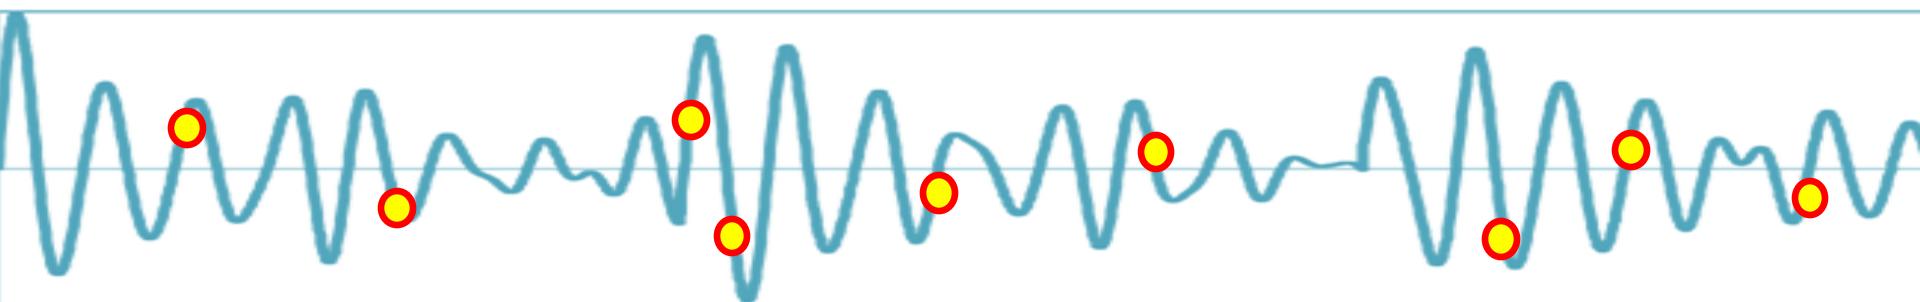
Compressed sensing continued

- Take a random projection of that sparse vector:

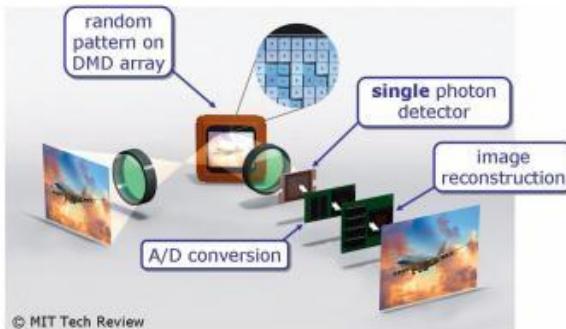


Why would I want to do that?

- Image compression and signal processing
- Especially when you **never have space to store the whole sparse vector to begin with.**



Randomly sampling (in the time domain) a signal that is sparse in the Fourier domain.

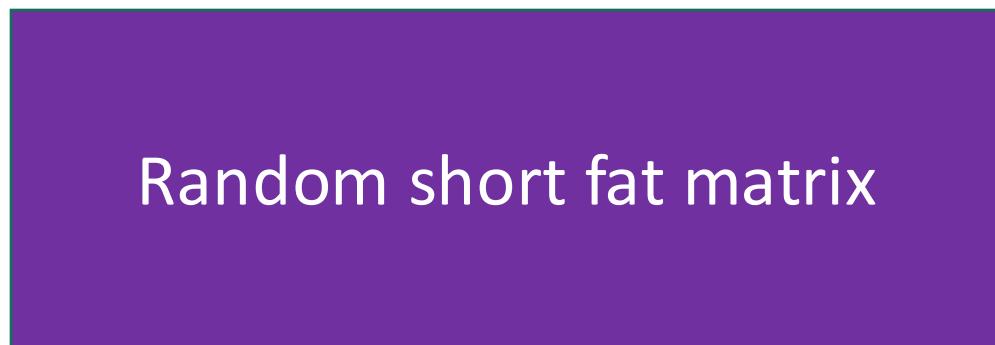


A “single pixel camera” is a thing.

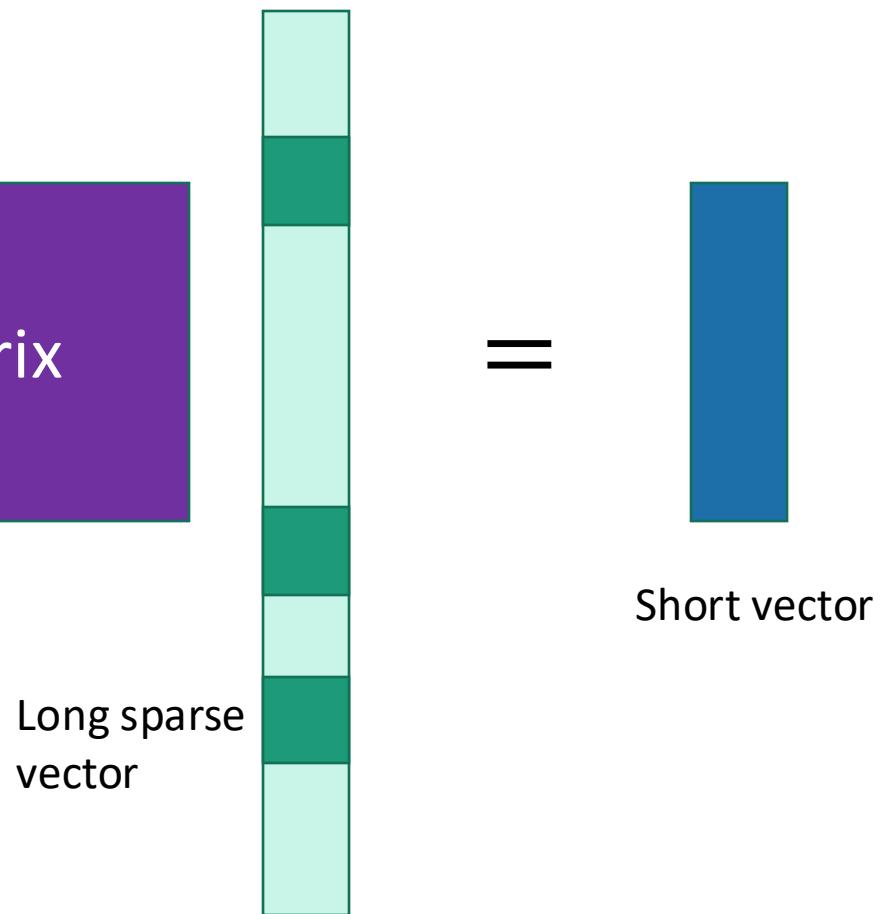
Random measurements in an fMRI means you spend less time inside an fMRI



All examples of this:



Goal: Given the **short vector**, recover the **long sparse vector**.



But why should this be possible?

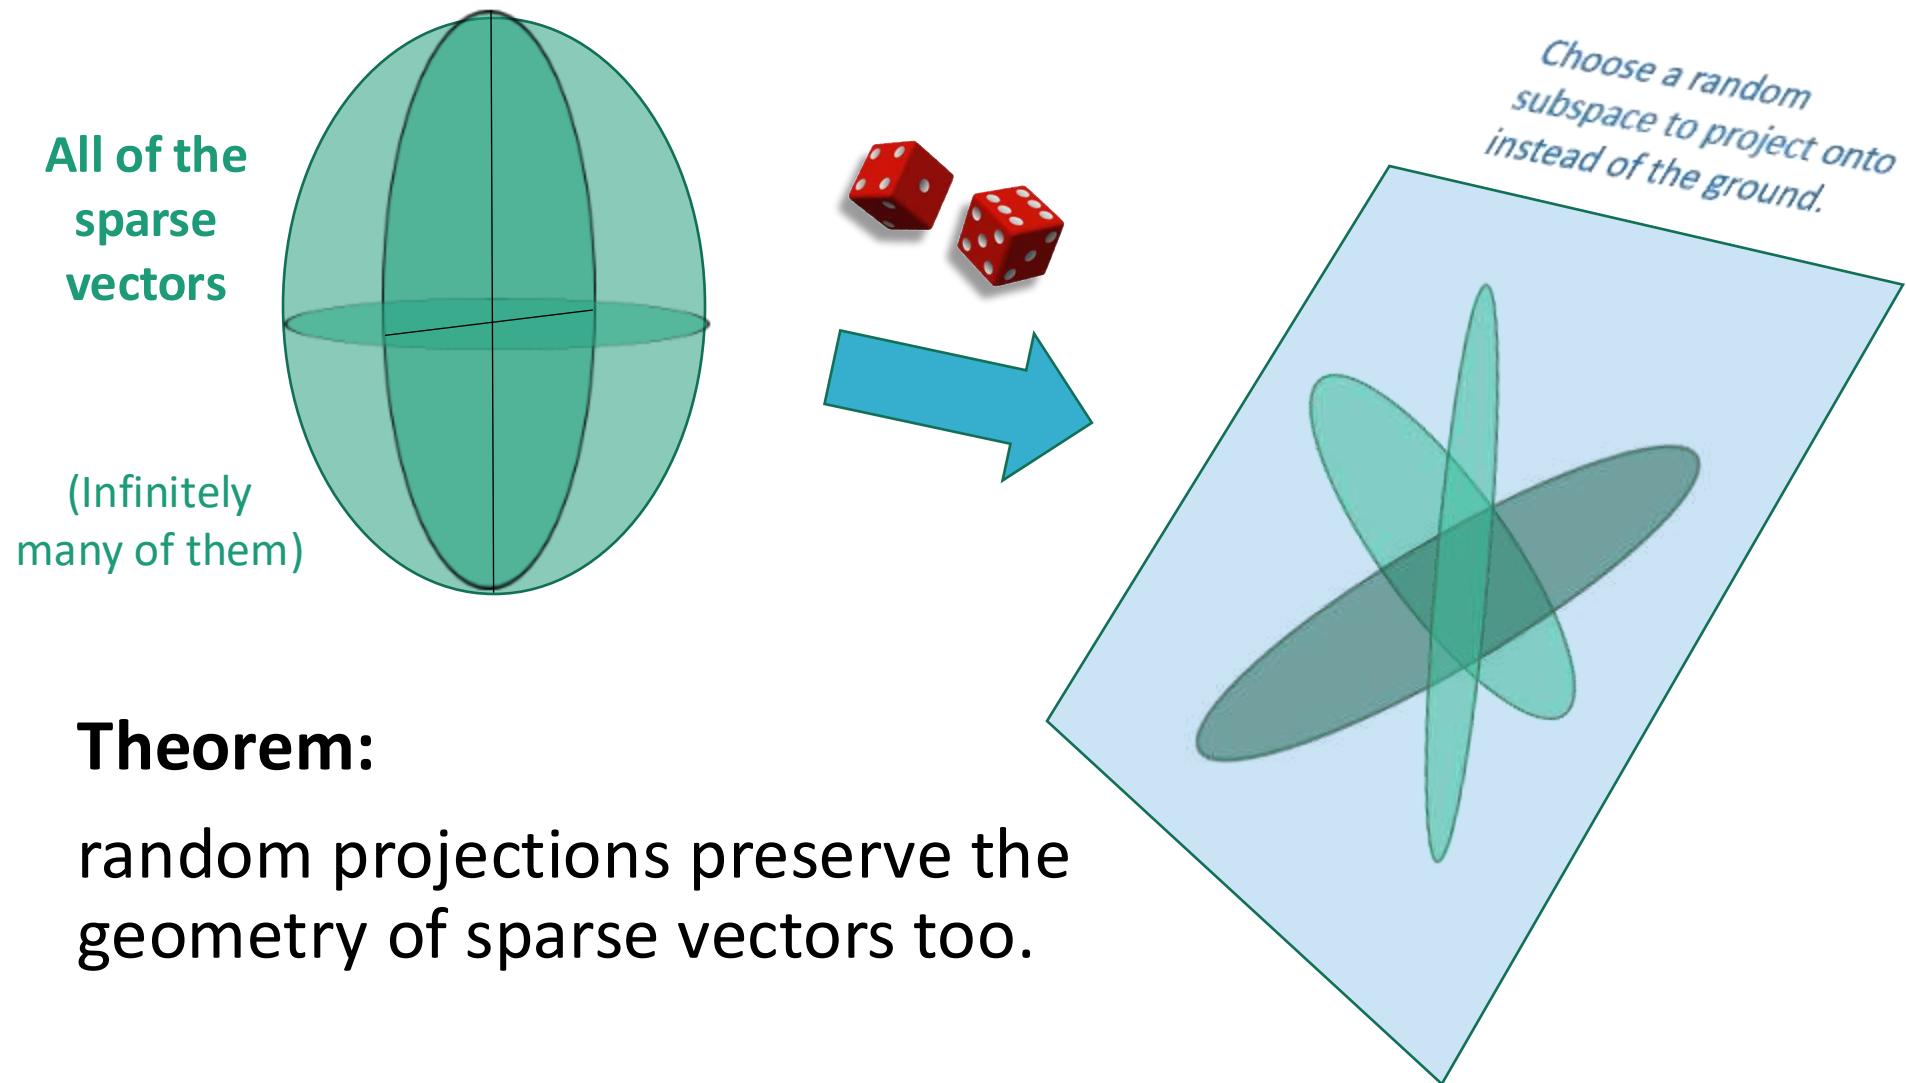
- There are tons of long vectors that map to the short vector!

Random short fat matrix

Goal: Given the **short vector**, recover the **long sparse vector.**

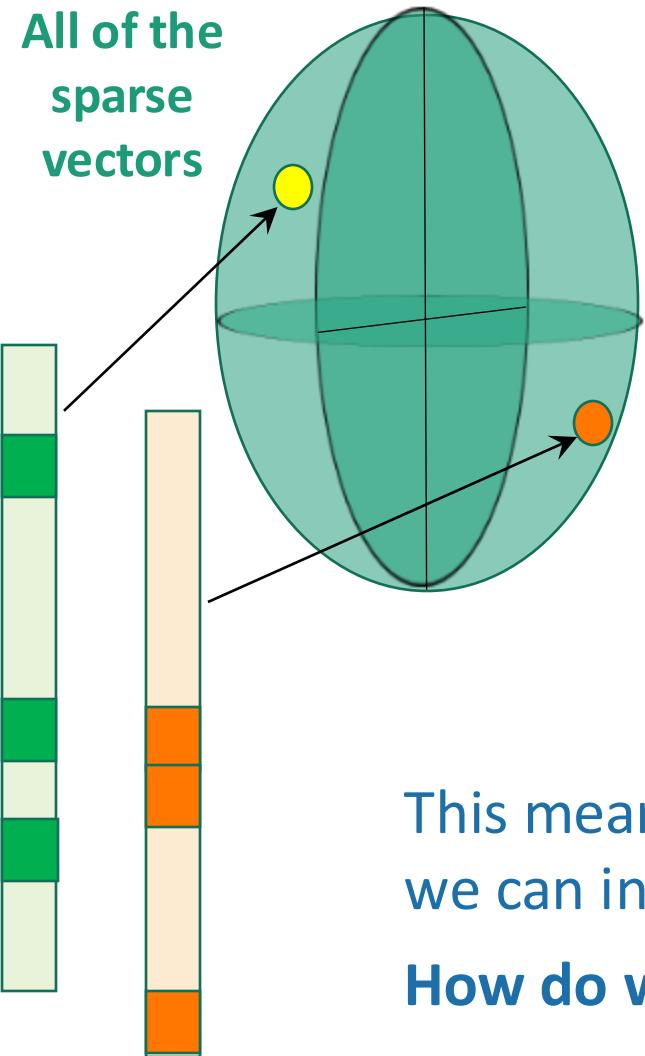
The diagram illustrates a mathematical operation involving a matrix and a vector. On the left, a large purple square labeled "Matrix" is multiplied by a tall, thin vertical vector. This vector is composed of several horizontal segments in different colors: light blue, dark green, light blue, dark green, light blue, dark green, and light blue. The result of this multiplication is an equals sign followed by a short, thick blue vertical bar labeled "Short vector".

Back to the geometry



We can (in theory) recover the big vector from the small one!

All of the sparse vectors

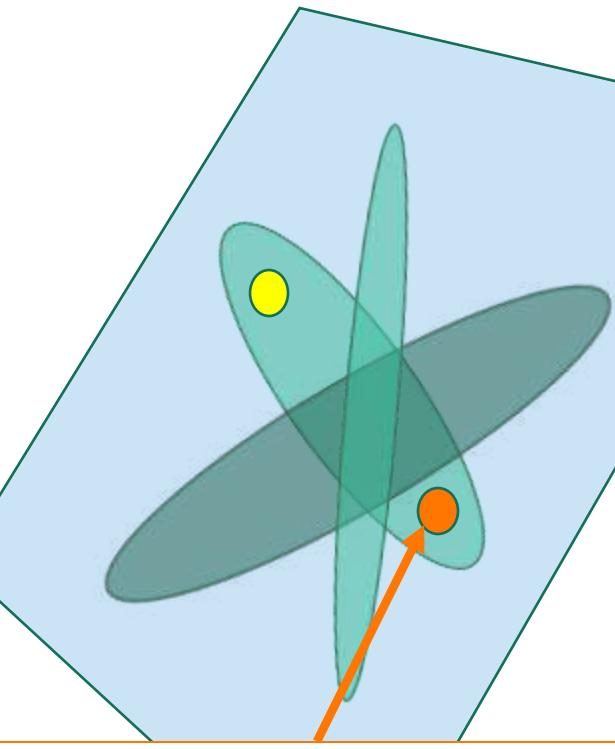


Multiply by

Random short
fat matrix

This means that, in theory,
we can invert that arrow.

How do we do this efficiently??



There may be tons of vectors
that map to this point, but only
one of them is sparse!

Goal: Given the **short vector**,
recover the **long sparse vector**.

An efficient algorithm?

What we'd like to do is:

Minimize number of
nonzero entries in x

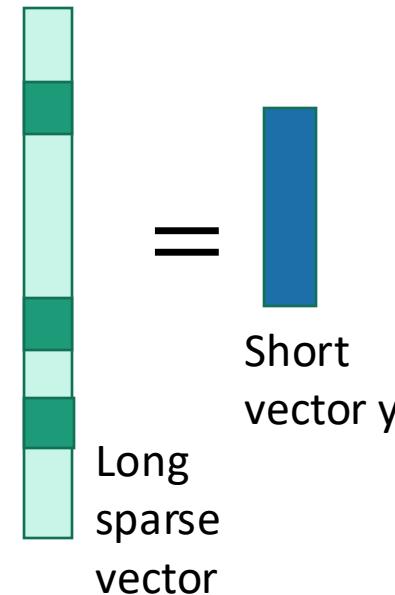
This norm is the sum
of the absolute values
of the entries of x

Instead:

Minimize $\|x\|_1$

Random short
fat matrix A

$$Ax = y$$



s.t.

Problem: I don't know
how to do that efficiently!

s.t.

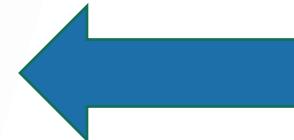
$$Ax = y$$

- It turns out that because the geometry of sparse vectors is preserved, this optimization problem **gives the same answer**.
- We can use **linear programming** to solve this quickly!

Today

A few gems

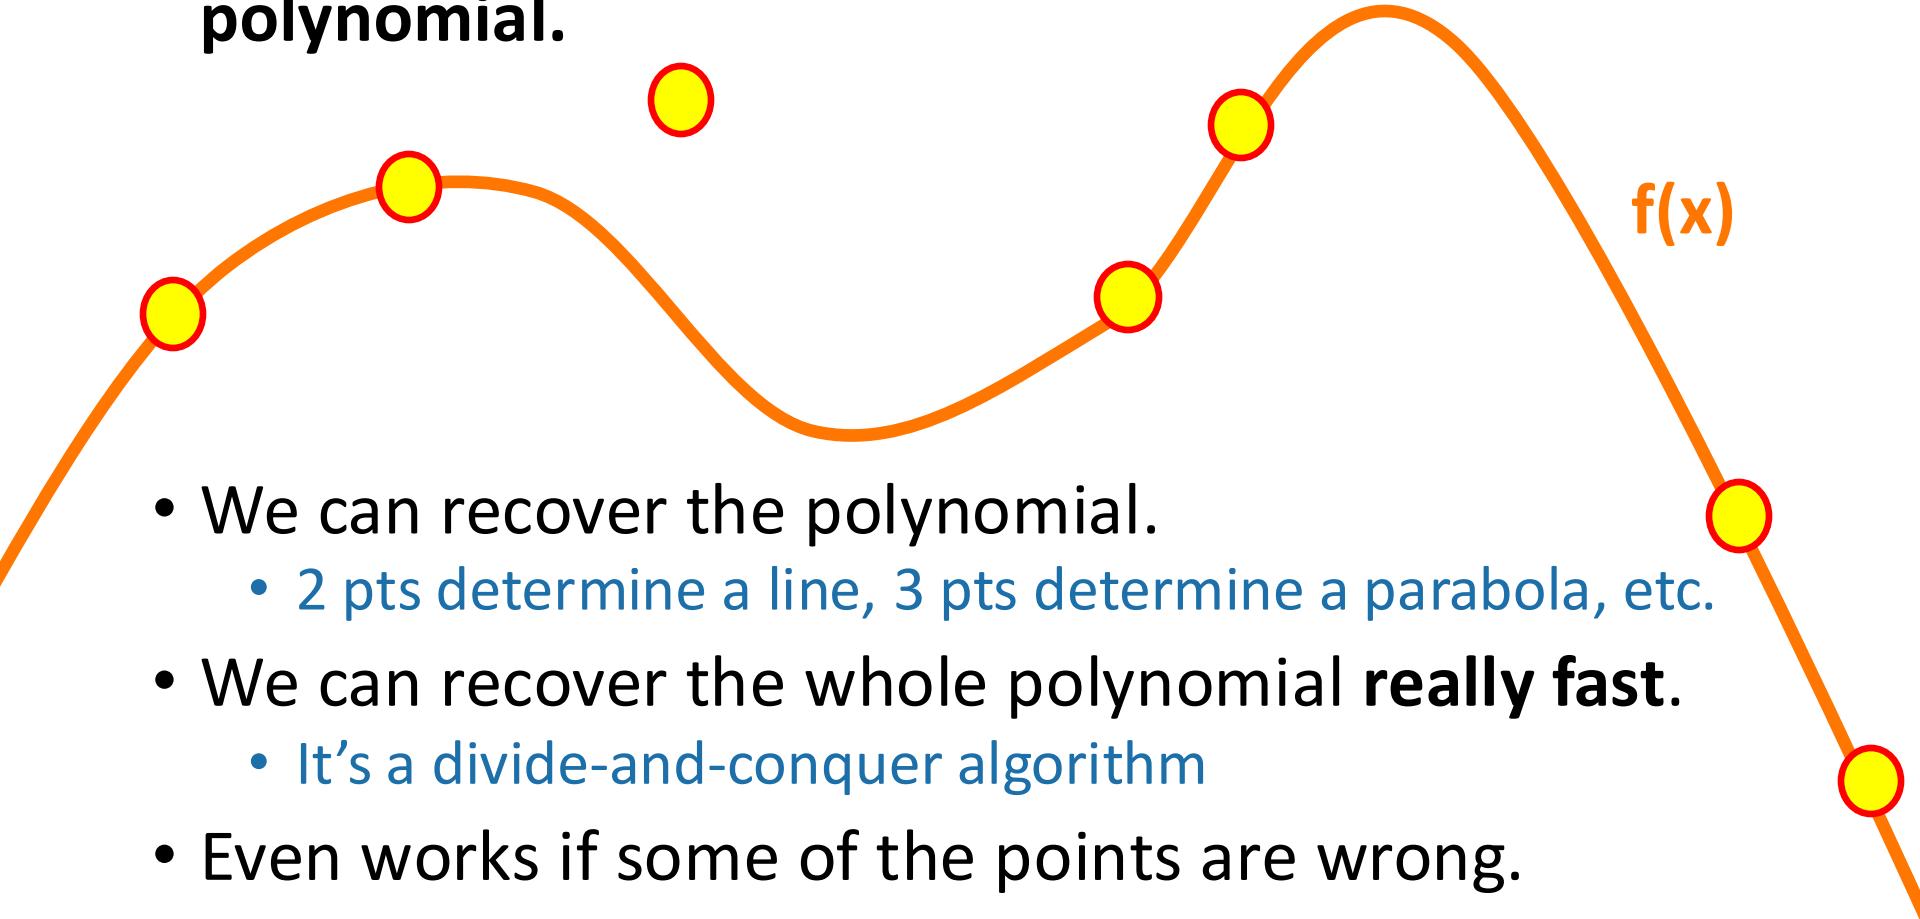
- Linear programming
- Random projections
- Low-degree polynomials



Another of my favorite tricks

Polynomial interpolation

- Say we have a few evaluation points of a **low-degree polynomial**.

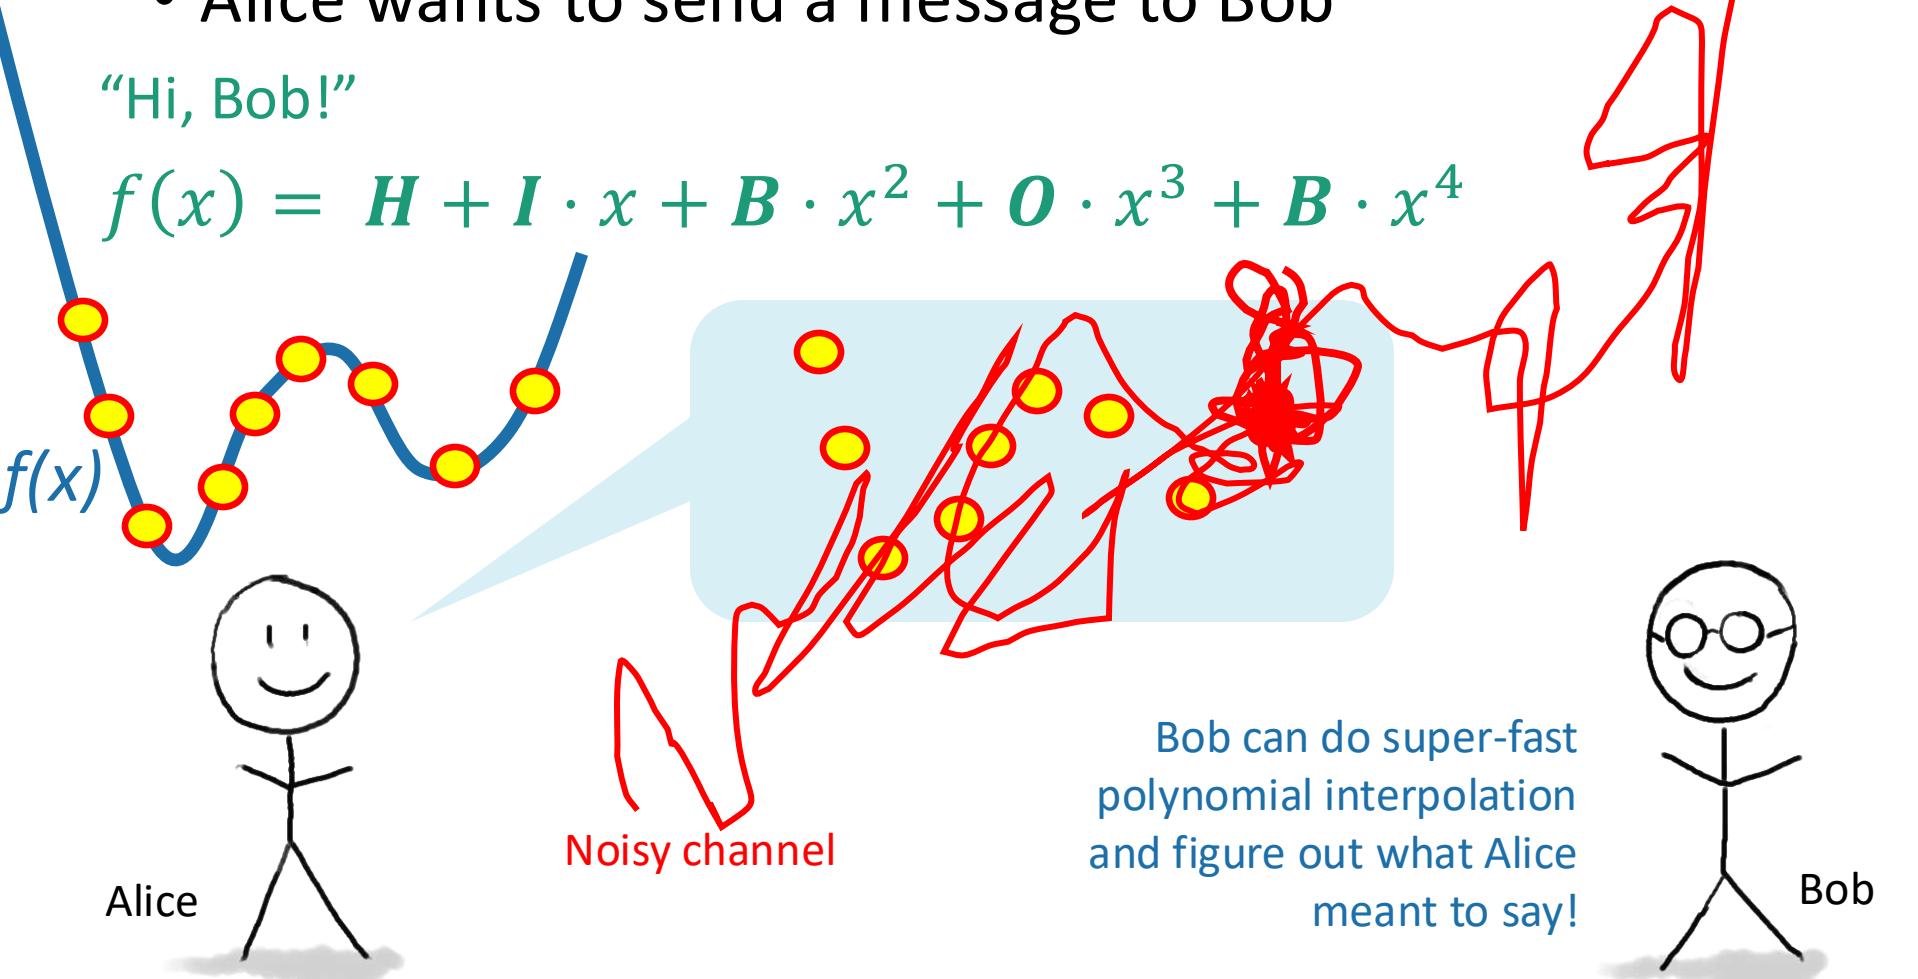


One application: Communication and Storage

- Alice wants to send a message to Bob

"Hi, Bob!"

$$f(x) = H + I \cdot x + B \cdot x^2 + O \cdot x^3 + B \cdot x^4$$



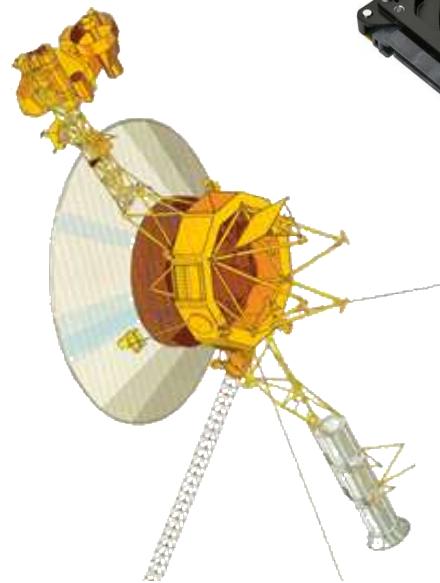
Bob can do super-fast
polynomial interpolation
and figure out what Alice
meant to say!

Alice

Bob

This is actually used in practice

- It's called “Reed-Solomon Encoding”

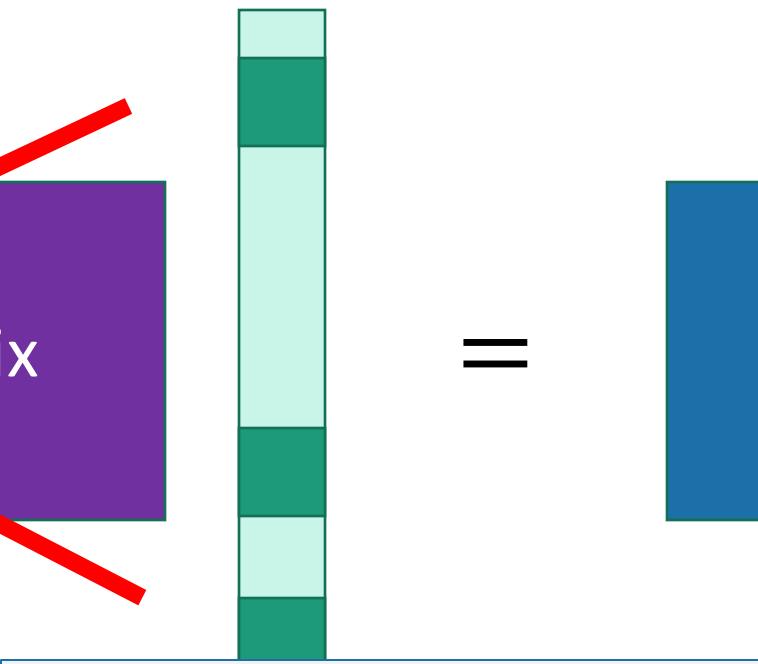


Another application:

Designing “random” projections that are better than random



The matrix that treats the big long vector as Alice's message polynomial and evaluates it REAL FAST at random points.



- This is still “random enough” to make the LP solution work.
- It is much more efficient to manipulate and store!

Today

A few gems

- Linear programming
- Random projections
- Low-degree polynomials



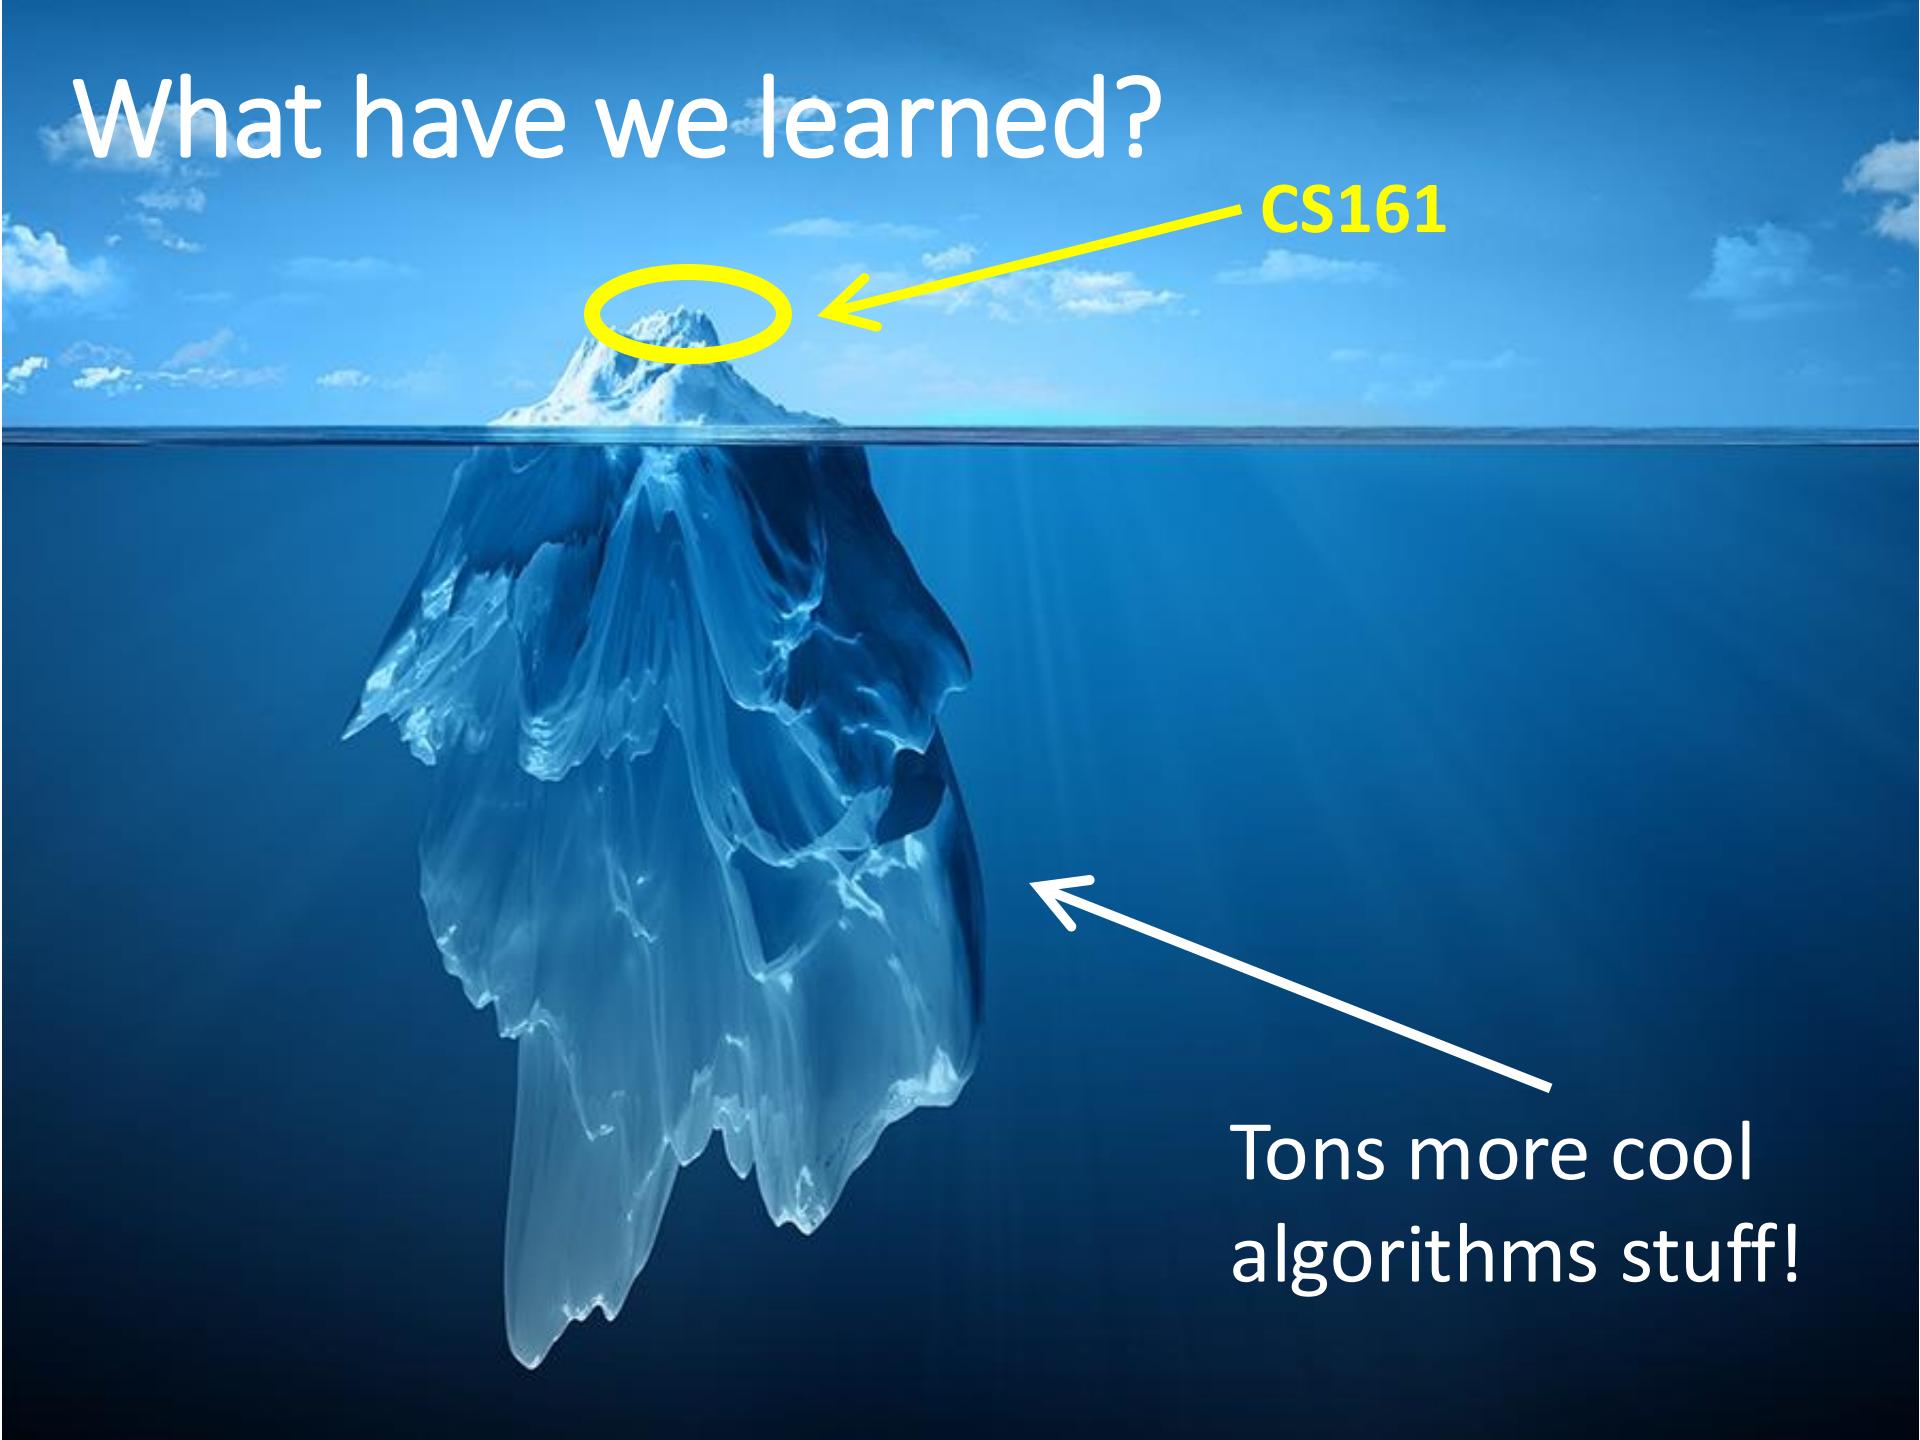
To learn more:

CS168, CS261, ...

CS168, CS261,
CS265, ...

CS168, CS250, ...

What have we learned?



CS161

Tons more cool
algorithms stuff!

To see more...

- Take more classes!
- Come hang out with the theory group!
 - Theory lunch, Thursdays at noon
 - Go to theory.stanford.edu and click on “Theory Lunch”

theory.stanford.edu

Stanford theory group:
We are very friendly.



A few final messages...

1. Some “Thank You”s...

- Thanks JP and the CGOE staff for helping out with recording all quarter long!

- Thanks to our course coordinator Amelie Byun!



- Thanks to our awesome embedded ethiCS team!



Justin



Jessica



Louie

2. Thanks to the TAs!!!

tell them you appreciate them!



Anisha



Isabel



Karan



Ly-Ly



Sophia



Ziyi



Will



Zayn



Maya



Ta-Wei



Spencer



Simon



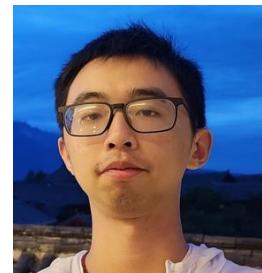
Bradley



Andy



Auddithio



Xiao



Mingwei



James

3.

THANKS
to you!!!!!!

